

Introduction

Natural Language Inferencing (NLI) task is one of the most important subsets of Natural Language Processing (NLP). NLI is also known as Recognizing Textual Entailment (RTE). NLI is a task of determining whether the given "hypothesis" and "premise" logically follow (entailment) or unfollow (contradiction) or are undetermined (neutral) to each other. In other words, the task of NLI is also defined as determining whether a "hypothesis" is true (entailment), false (contradiction), or undetermined (neutral) given a "premise". Table 1 gives some examples of these three cases.

Text (Premise)	Hypothesis	Judgement (Label)
A soccer game with multiple males playing. ---	Some men are playing a sport. ---	entailment ---
A black race car starts up in front of a crowd of people. ---	A man is driving down a lonely road. ---	contradiction ---
An older and a younger man smiling. ---	Two men are smiling and laughing at the cats playing on the floor. ---	neutral ---

Type *Markdown* and *LaTeX*: α^2

In [1]:

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import os
6 import re
7 from nltk.corpus import stopwords
8 import pickle
9 from tqdm import tqdm
10 from wordcloud import WordCloud
11 from sklearn.feature_extraction.text import TfidfVectorizer
12 import tensorflow as tf
13 import random
14 import math
15 from textblob import TextBlob
16 from collections import Counter
17 from sklearn.metrics.pairwise import cosine_similarity
18 from sklearn.feature_extraction.text import TfidfVectorizer
19 import nltk, string
20 from collections import Counter
21 from nltk.stem import WordNetLemmatizer
22 import Levenshtein
23 %matplotlib inline

```

In [2]: 1 nltk.download('punkt')

```
[nltk_data] Downloading package punkt to
[nltk_data]     C:\Users\anike\AppData\Roaming\nltk_data...
[nltk_data]     Package punkt is already up-to-date!
```

Out[2]: True

In [2]: 1 # Setting seed
2 global_seed = 42
3 np.random.seed(42)
4 random.seed(42)
5 tf.random.set_seed(42)

We have already extracted some features in the exploratory phase of the project. We will directly import data with those extracted features. We will start modelling with those features only. Then we will add features using word embeddings

Extracting word embedding features:average glove

In [11]: 1 embeddings_index = dict()
2 f = open('glove.6B.300d.txt', encoding="utf8")
3 for line in f:
4 values = line.split()
5 word = values[0]
6 coefs = np.asarray(values[1:], dtype='float32')
7 embeddings_index[word] = coefs
8 f.close()

In [12]: 1 glove_words=set(embeddings_index.keys())

train sentence1 sentence representations

In [13]:

```

1 from tqdm import tqdm
2 x_train_s1_tfidf2v = []
3 for sentence in tqdm(train_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_train_s1_tfidf2v.append(vector)
13
14 print(len(x_train_s1_tfidf2v))
15 print(len(x_train_s1_tfidf2v[0]))

```

100% | 5
49361/549361 [00:22<00:00, 23936.04it/s]

549361

300

train sentence2 sentence representations

In [14]:

```

1 from tqdm import tqdm
2 x_train_s2_tfidf2v = []
3 for sentence in tqdm(train_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_train_s2_tfidf2v.append(vector)
13
14 print(len(x_train_s2_tfidf2v))
15 print(len(x_train_s2_tfidf2v[0]))

```

100% | 5
49361/549361 [00:16<00:00, 33748.84it/s]

549361

300

validation sentence1 sentence representations

In [15]:

```

1 from tqdm import tqdm
2 x_val_s1_tfidf2v = []
3 for sentence in tqdm(val_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_val_s1_tfidf2v.append(vector)
13
14 print(len(x_val_s1_tfidf2v))
15 print(len(x_val_s1_tfidf2v[0]))

```

100% |██████████| 9842/9842 [00:00<00:00, 27434.12it/s]

9842

300

validation sentence2 sentence representations

In [16]:

```

1 from tqdm import tqdm
2 x_val_s2_tfidf2v = []
3 for sentence in tqdm(val_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_val_s2_tfidf2v.append(vector)
13
14 print(len(x_val_s2_tfidf2v))
15 print(len(x_val_s2_tfidf2v[0]))

```

100% |██████████| 9842/9842 [00:00<00:00, 44841.91it/s]

9842

300

test sentence1 sentence representations

In [17]:

```

1 from tqdm import tqdm
2 x_test_s1_tfidf2v = []
3 for sentence in tqdm(test_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_test_s1_tfidf2v.append(vector)
13
14 print(len(x_test_s1_tfidf2v))
15 print(len(x_test_s1_tfidf2v[0]))

```

100% | 9824/9824 [00:00<00:00, 28765.11it/s]

9824

300

test sentence2 sentence representations

In [18]:

```

1 from tqdm import tqdm
2 x_test_s2_tfidf2v = []
3 for sentence in tqdm(test_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         if word in glove_words:
8             vector += embeddings_index[word]
9             cnt_words += 1
10    if cnt_words != 0:
11        vector /= cnt_words
12    x_test_s2_tfidf2v.append(vector)
13
14 print(len(x_test_s2_tfidf2v))
15 print(len(x_test_s2_tfidf2v[0]))

```

100% | 9824/9824 [00:00<00:00, 40082.10it/s]

9824

300

```
In [19]: 1 print(type(x_train_s1_tfidf2v))
2 print(type(x_train_s1_tfidf2v[0]))
```

<class 'list'>
<class 'numpy.ndarray'>

```
In [20]: 1 print(type(X_train))
```

<class 'pandas.core.frame.DataFrame'>

```
In [21]: 1 x_train_s1_tfidf2v = np.vstack(x_train_s1_tfidf2v)
2 x_train_s2_tfidf2v = np.vstack(x_train_s2_tfidf2v)
3 x_val_s1_tfidf2v = np.vstack(x_val_s1_tfidf2v)
4 x_val_s2_tfidf2v = np.vstack(x_val_s2_tfidf2v)
5
```

```
In [22]: 1 x_test_s1_tfidf2v = np.vstack(x_test_s1_tfidf2v)
2 x_test_s2_tfidf2v = np.vstack(x_test_s2_tfidf2v)
```

```
In [23]: 1 np.save('x_train_s1_tfidf2v.npy', x_train_s1_tfidf2v)
2 np.save('x_train_s2_tfidf2v.npy', x_train_s2_tfidf2v)
3 np.save('x_val_s1_tfidf2v.npy', x_val_s1_tfidf2v)
4 np.save('x_val_s2_tfidf2v.npy', x_val_s2_tfidf2v)
5 np.save('x_test_s1_tfidf2v.npy', x_test_s1_tfidf2v)
6 np.save('x_test_s2_tfidf2v.npy', x_test_s2_tfidf2v)
```

```
In [13]: 1 x_train_s1_tfidf2v = np.load('x_train_s1_tfidf2v.npy')
2 x_train_s2_tfidf2v = np.load('x_train_s2_tfidf2v.npy')
3 x_val_s1_tfidf2v = np.load('x_val_s1_tfidf2v.npy')
4 x_val_s2_tfidf2v = np.load('x_val_s2_tfidf2v.npy')
5 x_test_s1_tfidf2v = np.load('x_test_s1_tfidf2v.npy')
6 x_test_s2_tfidf2v = np.load('x_test_s2_tfidf2v.npy')
```

cosine similarity between average glove

```
In [29]: 1 def tfidf_w2v_cosine(mat1, mat2):
2     """
3         Given two matrices, returns a list of cosine similarities between rows of
4     """
5
6     from scipy.spatial.distance import cosine
7     sim = [1-cosine(i,j) for i,j in zip(mat1, mat2)]
8     return sim
9
10
11
```

```
In [30]: 1 train_tfidf_w2v_cosine = tfidf_w2v_cosine(x_train_s1_tfidf2v, x_train_s2_tf
C:\Users\anike\conda\envs\tf_test\lib\site-packages\scipy\spatial\distance.py:
720: RuntimeWarning: invalid value encountered in double_scalars
    dist = 1.0 - uv / np.sqrt(uu * vv)
```

```
In [36]: 1 val_tfidf_w2v_cosine = tfidf_w2v_cosine(x_val_s1_tfidf2v, x_val_s2_tfidf2v)
```

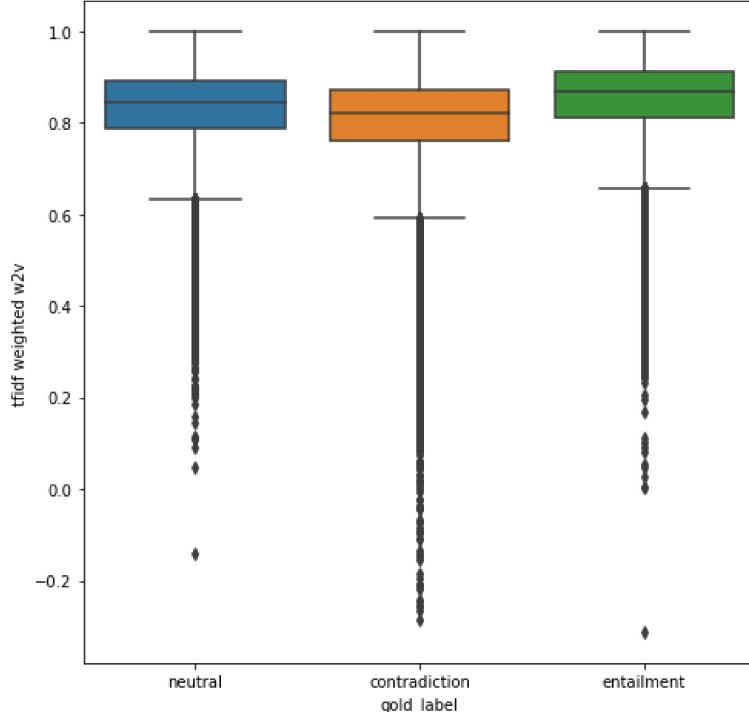
```
In [41]: 1 test_tfidf_w2v_cosine = tfidf_w2v_cosine(x_test_s1_tfidf2v, x_test_s2_tfidf2v)
```

```
In [42]: 1 train_df['tfidf_w2v_cosine'] = train_tfidf_w2v_cosine
2 val_df['tfidf_w2v_cosine'] = val_tfidf_w2v_cosine
3 test_df['tfidf_w2v_cosine'] = test_tfidf_w2v_cosine
```

```
In [69]: 1 train_df = train_df.fillna(0)
```

```
In [43]: 1 fig,ax = plt.subplots(1,1, figsize=(7,7))
2 sns.boxplot(x = train_df['gold_label'], y = train_df['tfidf_w2v_cosine'], ax
3 ax.set_title('box plot of tfidf weighted w2v between premise and hypothesis')
4 ax.set_ylabel('tfidf weighted w2v')
5 plt.tight_layout()
6 plt.show()
```

box plot of tfidf weighted w2v between premise and hypothesis for each label



As expected, the tfidf weighted w2v similarity between two sentences is more for entailment sentences. It is least for contradictory sentences

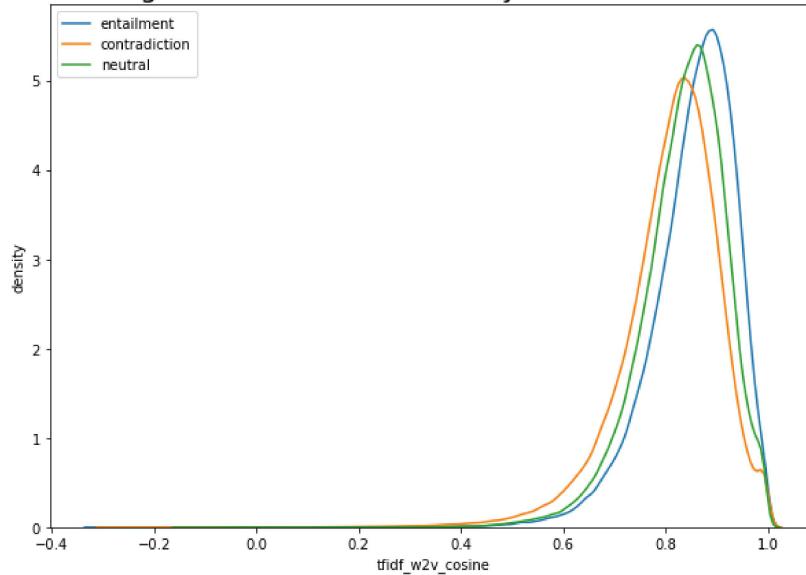
In [44]:

```

1 fig,ax = plt.subplots(1,1, figsize=(10,7))
2 sns.kdeplot(x = train_df[train_df['gold_label'] == 'entailment']['tfidf_w2v_'
3 sns.kdeplot(x = train_df[train_df['gold_label'] == 'contradiction']['tfidf_w'
4 sns.kdeplot(x = train_df[train_df['gold_label'] == 'neutral']['tfidf_w2v_cos'
5 ax.set_xlabel('tfidf_w2v_cosine')
6 ax.set_ylabel('density')
7 ax.set_title('distribution of tfidf weighted w2v cosine similarity between'
8 plt.legend()
9 plt.show()

```

distribution of tfidf weighted w2v cosine similarity between sentence 1 and sentence 2



In [70]:

```

1 train_df.to_csv('train_df.csv')
2 val_df.to_csv('val_df.csv')
3 test_df.to_csv('test_df.csv')

```

In [71]:

```

1 train_df = pd.read_csv('train_df.csv', index_col = 0)
2 val_df = pd.read_csv('val_df.csv', index_col = 0)
3 test_df = pd.read_csv('test_df.csv', index_col = 0)

```

In [72]:

```

1 to_keep_x = ['sentence1_num_words', 'sentence2_num_words', 'log_bleu_score',
2             'cosine_similarity', 'tfidf_cosine_sim', 'diff_num_words', 's1_CC',
3             's1_IN', 's1_JJ', 's1_JJR', 's1_JJS', 's1_PRP', 's1_PRP$', 's1_RB',
4             's2_JJ', 's2_JJR', 's2_JJS', 's2_NN', 's2_NNP', 's2_NNS', 's2_PRP',
5             's2_PRP$', 's2_RB', 's2_RBR', 's2_RBS', 's2_UH', 's2_VB', 's2_VBD',
6             's2_VBG', 's2_VBN', 's2_VBP', 's2_VBZ', 'ratio_noun', 'ratio_verb', '
7
8 X_train = train_df[to_keep_x]
9 y_train = train_df['gold_label']
10 X_val = val_df[to_keep_x]
11 y_val = val_df['gold_label']

```

Training random forest model without average

glove vectors

```
In [73]: 1 from sklearn.ensemble import RandomForestClassifier  
2 model2 = RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_sample  
3                                     n_estimators=800)  
4 model2.fit(X_train, y_train)
```

```
Out[73]: RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_samples_split=10,  
                                 n_estimators=800)
```

```
In [75]: 1 val_pred = model2.predict(X_val)  
2  
3 from sklearn.metrics import accuracy_score  
4 acc = accuracy_score(y_val, val_pred)  
5  
6 acc
```

```
Out[75]: 0.5821987400934769
```

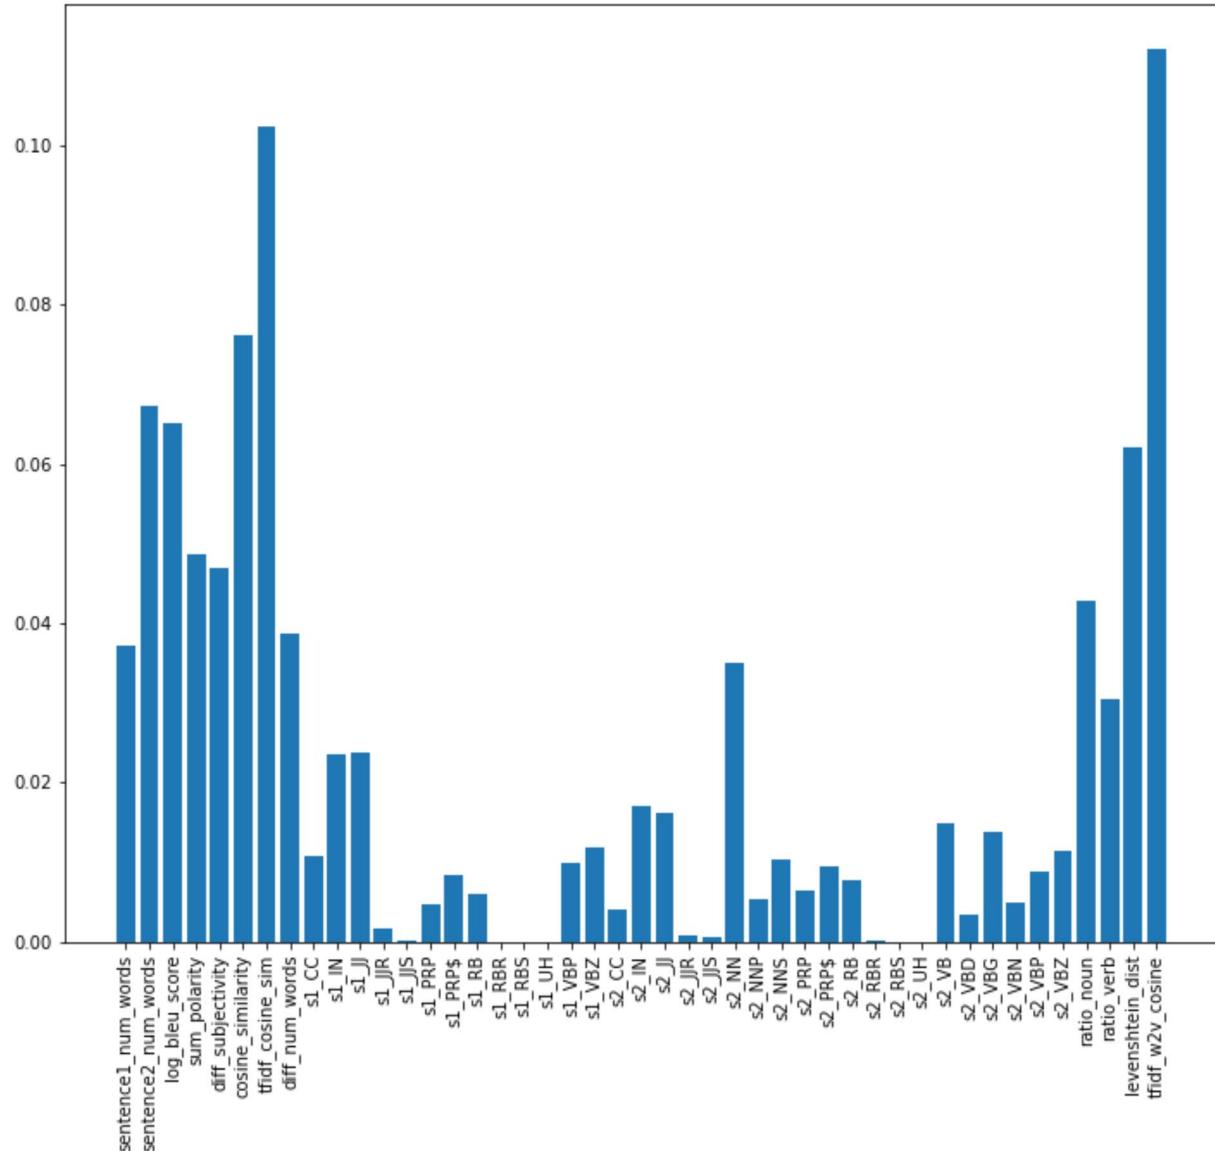
with only extracted features we are getting accuracy of 0.58

In [77]:

```

1 importance = model2.feature_importances_
2
3 # Plotting feature importance
4 columns = train_df[to_keep_x].columns
5 x_pos = np.arange(len(columns))
6 plt.figure(figsize = (12,10))
7 plt.bar(x_pos, importance)
8 plt.xticks(x_pos, columns, rotation='vertical')
9 plt.show()
10 #plt.plot(np.array(range(0,47)), importance, 'bo' )

```



Training model with extracted features + avg glove vectors of both sentences

In [84]:

```

1 x_train = np.hstack((X_train, x_train_s1_tfidf2v, x_train_s2_tfidf2v))
2 x_val = np.hstack((X_val, x_val_s1_tfidf2v, x_val_s2_tfidf2v))

```

```
In [86]: 1 from sklearn.ensemble import RandomForestClassifier
2 model3 = RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_samples_split=10,
3                                 n_estimators=800)
4 model3.fit(x_train, y_train)
```

Out[86]: RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_samples_split=10, n_estimators=800)

```
In [88]: 1 val_pred = model3.predict(x_val)
2
3 from sklearn.metrics import accuracy_score
4 acc = accuracy_score(y_val, val_pred)
5
6 acc
```

Out[88]: 0.6633814265393213

after including average word vectors accuracy increases to 0.66

Considering ngrams: tfidf cosine similarity with bigrams, trigrams and quadgrams

```
In [89]: 1 #https://stackoverflow.com/questions/8897593/how-to-compute-the-similarity-between-strings
2 stemmer = nltk.stem.porter.PorterStemmer()
3 remove_punctuation_map = dict((ord(char), None) for char in string.punctuation)
4
5 def stem_tokens(tokens):
6     return [stemmer.stem(item) for item in tokens]
7
8 '''remove punctuation, lowercase, stem'''
9 def normalize(text):
10     return stem_tokens(nltk.word_tokenize(text.lower()).translate(remove_punctuation_map))
11
12 vectorizer = TfidfVectorizer(tokenizer=normalize, ngram_range=(2, 4))#we gave it a name
13
14 def cosine_sim(text1, text2):
15     tfidf = vectorizer.fit_transform([text1, text2])
16     return ((tfidf * tfidf.T).A)[0,1]
```

ngram tfidf cosine similarity for training data

```
In [90]: 1 tfidf_cos_sim_ngram = list()
2 for i, row in train_df.iterrows():
3     cos_sim = cosine_sim(row['sentence1'], row['sentence2'])
4     tfidf_cos_sim_ngram.append(cos_sim)
5
6 train_df['tfidf_cosine_sim_ngram'] = tfidf_cos_sim_ngram
```

ngram tfidf cosine similarity for val data

In [91]:

```

1 tfidf_cos_sim_ngram = list()
2 for i,row in val_df.iterrows():
3     cos_sim = cosine_sim(row['sentence1'], row['sentence2'])
4     tfidf_cos_sim_ngram.append(cos_sim)
5
6 val_df['tfidf_cosine_sim_ngram'] = tfidf_cos_sim_ngram

```

ngram tfidf cosine similarity for test data

In [92]:

```

1 tfidf_cos_sim_ngram = list()
2 for i,row in test_df.iterrows():
3     cos_sim = cosine_sim(row['sentence1'], row['sentence2'])
4     tfidf_cos_sim_ngram.append(cos_sim)
5
6 test_df['tfidf_cosine_sim_ngram'] = tfidf_cos_sim_ngram

```

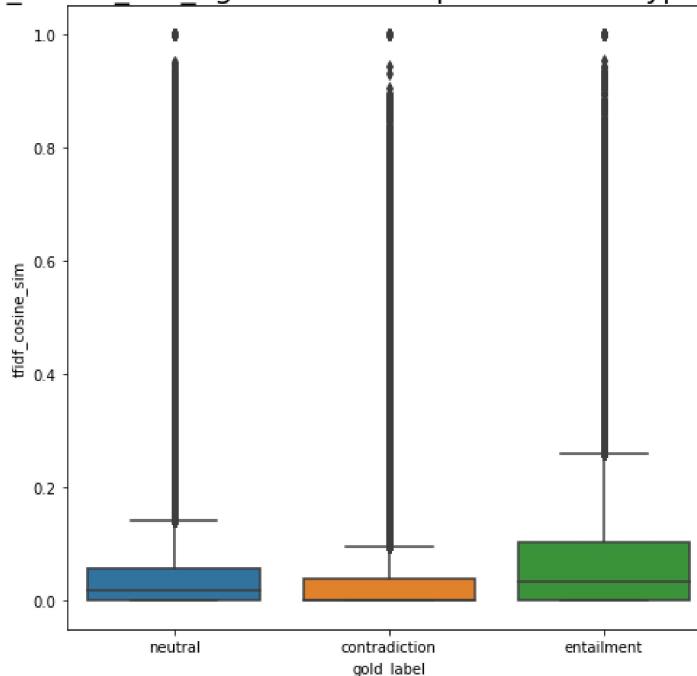
In [94]:

```

1 fig,ax = plt.subplots(1,1, figsize=(7,7))
2 sns.boxplot(x = train_df['gold_label'], y = train_df['tfidf_cosine_sim_ngram']
3 ax.set_title('box plot of tfidf_cosine_sim_ngram between premise and hypothesis')
4 ax.set_ylabel('tfidf_cosine_sim')
5 plt.tight_layout()
6 plt.show()
7
8

```

box plot of tfidf_cosine_sim_ngram between premise and hypothesis for each label

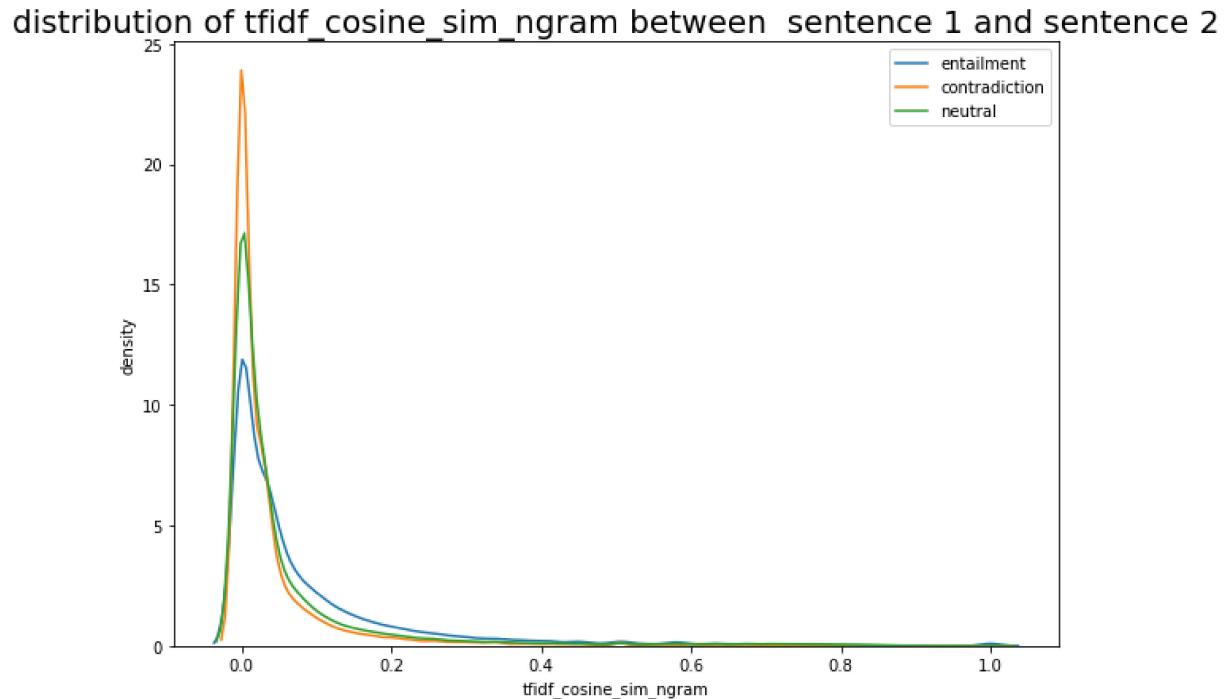


In [95]:

```

1 fig,ax = plt.subplots(1,1, figsize=(10,7))
2 sns.kdeplot(x = train_df[train_df['gold_label'] == 'entailment']['tfidf_cosine_sim_ngram'])
3 sns.kdeplot(x = train_df[train_df['gold_label'] == 'contradiction']['tfidf_cosine_sim_ngram'])
4 sns.kdeplot(x = train_df[train_df['gold_label'] == 'neutral']['tfidf_cosine_sim_ngram'])
5 ax.set_xlabel('tfidf_cosine_sim_ngram')
6 ax.set_ylabel('density')
7 ax.set_title('distribution of tfidf_cosine_sim_ngram between sentence 1 and sentence 2')
8 plt.legend()
9 plt.show()

```



In [99]:

```

1 train_df.to_csv('train_df.csv')
2 val_df.to_csv('val_df.csv')
3 test_df.to_csv('test_df.csv')

```

In [5]:

```

1 train_df = pd.read_csv('train_df.csv', index_col = 0)
2 val_df = pd.read_csv('val_df.csv', index_col = 0)
3 test_df = pd.read_csv('test_df.csv', index_col = 0)

```

In [3]:

```

1 to_keep = ['sentence1_num_words', 'sentence2_num_words',
2          'log_bleu_score', 'sentence1_polarity',
3          'sentence2_polarity', 'sentence1_subjectivity',
4          'sentence2_subjectivity',
5          'cosine_similarity', 'tfidf_cosine_sim', 's1_CC',
6          's1_IN', 's1_JJ', 's1_JJR', 's1_JJS', 's1_NN', 's1_NNP', 's1_NNS',
7          's1_PRP', 's1_PRP$', 's1_RB', 's1_RBR', 's1_RBS', 's1_UH', 's1_VB',
8          's1_VBD', 's1_VBG', 's1_VBN', 's1_VBP', 's1_VBZ', 's2_CC', 's2_IN',
9          's2_JJ', 's2_JJR', 's2_JJS', 's2_NN', 's2_NNP', 's2_NNS', 's2_PRP',
10         's2_PRP$', 's2_RB', 's2_RBR', 's2_RBS', 's2_UH', 's2_VB', 's2_VBD',
11         's2_VBG', 's2_VBN', 's2_VBP', 's2_VBZ', 'levenshtein_dist', 'tfidf_w2v',
12         'tfidf_cosine_sim_ngram']

```

```
In [4]: 1 X_train = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val = val_df[to_keep]
4 y_val = val_df['gold_label']
```

```
In [7]: 1 x_train = np.hstack((X_train, x_train_s1_tfidf2v, x_train_s2_tfidf2v))
2 x_val = np.hstack((X_val, x_val_s1_tfidf2v, x_val_s2_tfidf2v))
```

```
In [26]: 1 from sklearn.ensemble import RandomForestClassifier
2 model4 = RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_samples_split=10,
3                                n_estimators=800)
4 model4.fit(x_train, y_train)
5
6
```

Out[26]: RandomForestClassifier(max_depth=80, min_samples_leaf=2, min_samples_split=10, n_estimators=800)

```
In [27]: 1 val_pred = model4.predict(x_val)
2 from sklearn.metrics import accuracy_score
3 acc = accuracy_score(y_val, val_pred)
4 acc
```

Out[27]: 0.662466978256452

random forest best score 0.66 on validation set

XGBoost

```
In [32]: 1 to_keep = ['log_bleu_score', 'sum_polarity', 'diff_subjectivity',
2             'cosine_similarity', 'tfidf_cosine_sim', 'diff_num_words', 's1_CC',
3             's1_IN', 's1_PRP', 's1_PRP$', 's1_RB', 's1_RBR', 's1_RBS', 's1_UH', 's1_UHS',
4             's2_PRP$', 's2_RB', 's2_RBR', 's2_RBS', 's2_UH', 'ratio_noun', 'ratio_nouns',
5             'diff_adjective', 'levenshtein_dist', 'tfidf_w2v_cosine',
6             'tfidf_cosine_sim_ngram']
```

```
In [5]: 1 X_train = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val = val_df[to_keep]
4 y_val = val_df['gold_label']
5
```

```
In [ ]: 1 x_train = np.hstack((X_train, x_train_s1_tfidf2v, x_train_s2_tfidf2v))
2 x_val = np.hstack((X_val, x_val_s1_tfidf2v, x_val_s2_tfidf2v))
```

```
In [7]: 1 #classes for grid search and cross-validation, function for splitting data a
2 from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, train_test_split
3 from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, confusion_matrix
4 X_tr, X_te, y_tr, y_te = train_test_split(X_train, y_train, test_size = 0.4,
```

In [10]:

```

1 # hyperparameter tuning
2 #XGBoost Library
3 import xgboost as xgb
4 from sklearn.model_selection import RandomizedSearchCV
5 param_grid = {'gamma': [0,0.1,0.4,0.8,3.2,6.4,25.6,102.4, 200],
6               'learning_rate': [0.01,0.001, 0.1, 0.5],
7               'max_depth': [3,4,5,9,12],
8               'n_estimators': [50,100,200,400,800],
9               'reg_alpha': [0,0.1,0.001,1,5,100,200],
10              'reg_lambda': [0,0.1,0.001,1,5,100,200]}
11 xgbc = xgb.XGBClassifier()
12 # Random search of parameters, using 3 fold cross validation,
13 # search across 100 different combinations, and use all available cores
14 xgb_random = RandomizedSearchCV(estimator = xgbc, param_distributions = para
15 # Fit the random search model
16 xgb_random.fit(X_tr, y_tr)

```

C:\Users\anike\.conda\envs\tf_test\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_classes - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[14:03:35] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Out[10]: RandomizedSearchCV(cv=4,

```

estimator=XGBClassifier(base_score=None, booster=None,
                       colsample_bytree=None,
                       colsample_bynode=None,
                       colsample_bylevel=None,
                       enable_categorical=False, gamma=None,
                       gpu_id=None, importance_type=None,
                       interaction_constraints=None,
                       learning_rate=None,
                       max_delta_step=None, max_depth=None,
                       min_child_weight=None, missing=nan,
                       monotone_constraints...,
                       subsample=None, tree_method=None,
                       validate_parameters=None,
                       verbosity=None),
n_iter=50, n_jobs=6,
param_distributions={'gamma': [0, 0.1, 0.4, 0.8, 3.2, 6.4,
                               25.6, 102.4, 200],
                     'learning_rate': [0.01, 0.001, 0.1, 0.5],
                     'max_depth': [3, 4, 5, 9, 12],
                     'n_estimators': [50, 100, 200, 400, 800],
                     'reg_alpha': [0, 0.1, 0.001, 1, 5, 100, 200],
                     'reg_lambda': [0, 0.1, 0.001, 1, 5, 100, 200]},

```

```
'reg_lambda': [0, 0.1, 0.001, 1, 5, 10
0,
200]},
random_state=42)
```

In [11]: 1 xgb_random.best_estimator_

Out[11]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, enable_categorical=False, gamma=0.1, gpu_id=-1, importance_type=None, interaction_constraints='', learning_rate=0.1, max_delta_step=0, max_depth=12, min_child_weight=1, missing=np.nan, monotone_constraints='()', n_estimators=200, n_jobs=16, num_parallel_tree=1, objective='multi:softprob', predictor='auto', random_state=0, reg_alpha=0, reg_lambda=100, scale_pos_weight=None, subsample=1, tree_method='exact', validate_parameters=1, verbosity=None)

In [14]: 1 X_train = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val = val_df[to_keep]
4 y_val = val_df['gold_label']
5 x_train = np.hstack((X_train, x_train_s1_tfidf2v, x_train_s2_tfidf2v))
6 x_val = np.hstack((X_val, x_val_s1_tfidf2v, x_val_s2_tfidf2v))

In [18]: 1 xgbc = xgb.XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, enable_categorical=False, gamma=0.1, gpu_id=-1, importance_type=None, interaction_constraints='', learning_rate=0.1, max_delta_step=0, max_depth=12, min_child_weight=1, missing=np.nan, monotone_constraints='()', n_estimators=200, n_jobs=16, num_parallel_tree=1, objective='multi:softprob', predictor='auto', random_state=0, reg_alpha=0, reg_lambda=100, scale_pos_weight=None, subsample=1, tree_method='exact', validate_parameters=1, verbosity=None)

In [22]: 1 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'], value=0)
2 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'], value=0)

In [25]: 1 print(x_train.shape)
2 print(x_val.shape)

```
(549361, 629)
(9842, 629)
```

In [23]: 1 xgbc.fit(x_train, y_train)

```
[14:15:04] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.
1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric
used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'.
Explicitly set eval_metric if you'd like to restore the old behavior.
```

Out[23]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
gamma=0.1, gpu_id=-1, importance_type=None,
interaction_constraints='', learning_rate=0.1, max_delta_step=0,
max_depth=12, min_child_weight=1, missing=nan,
monotone_constraints='()', n_estimators=200, n_jobs=16,
num_parallel_tree=1, objective='multi:softprob', predictor='auto',
random_state=0, reg_alpha=0, reg_lambda=100,
scale_pos_weight=None, subsample=1, tree_method='exact',
validate_parameters=1, verbosity=None)

In [26]: 1 val_pred = xgbc.predict(x_val)
2 from sklearn.metrics import accuracy_score
3 acc = accuracy_score(y_val, val_pred)
4 acc

Out[26]: 0.7087990245884983

xgboost best score 0.7 on validation set(with extracted features+average word embeddings(GLoVe))

Embeddings using fasttext

In [4]: 1 import fasttext

In [20]:

```

1 #Loading a retrained fasttext model
2 #https://fasttext.cc/docs/en/crawl-vectors.html
3 import fasttext.util
4 fasttext.util.download_model('en', if_exists='ignore') # English
5 ft = fasttext.load_model('cc.en.300.bin')

```

Downloading <https://dl.fbaipublicfiles.com/fasttext/vectors-crawl/cc.en.300.bin.gz> (<https://dl.fbaipublicfiles.com/fasttext/vectors-crawl/cc.en.300.bin.gz>)

train sentence1 average fasttext

In [23]:

```

1 from tqdm import tqdm
2 x_train_s1_avg_fasttext = []
3 for sentence in tqdm(train_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_train_s1_avg_fasttext.append(vector)
12
13 print(len(x_train_s1_avg_fasttext))
14 print(len(x_train_s1_avg_fasttext[0]))

```

100% |██████████|
549361/549361 [01:04<00:00, 8521.67it/s]

549361
300

train sentence2 average fasttext

In [24]:

```

1 from tqdm import tqdm
2 x_train_s2_avg_fasttext = []
3 for sentence in tqdm(train_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_train_s2_avg_fasttext.append(vector)
12
13 print(len(x_train_s2_avg_fasttext))
14 print(len(x_train_s2_avg_fasttext[0]))

```

100% |██████████| 5
49361/549361 [00:40<00:00, 13669.84it/s]

549361

300

val sentence1 average fasttext

In [25]:

```

1 from tqdm import tqdm
2 x_val_s1_avg_fasttext = []
3 for sentence in tqdm(val_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_val_s1_avg_fasttext.append(vector)
12
13 print(len(x_val_s1_avg_fasttext))
14 print(len(x_val_s1_avg_fasttext[0]))

```

100% |██████████| 5
9842/9842 [00:01<00:00, 7761.10it/s]

9842

300

val sentence2 average fasttext

In [26]:

```

1 from tqdm import tqdm
2 x_val_s2_avg_fasttext = []
3 for sentence in tqdm(val_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_val_s2_avg_fasttext.append(vector)
12
13 print(len(x_val_s2_avg_fasttext))
14 print(len(x_val_s2_avg_fasttext[0]))

```

100% |██████████| 9842/9842 [00:00<00:00, 12560.89it/s]

9842

300

test sentence1 average fasttext

In [27]:

```

1 from tqdm import tqdm
2 x_test_s1_avg_fasttext = []
3 for sentence in tqdm(test_df['sentence1_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_test_s1_avg_fasttext.append(vector)
12
13 print(len(x_test_s1_avg_fasttext))
14 print(len(x_test_s1_avg_fasttext[0]))

```

100% |██████████| 9824/9824 [00:01<00:00, 7749.76it/s]

9824

300

test sentence2 average fasttext

In [28]:

```

1 from tqdm import tqdm
2 x_test_s2_avg_fasttext = []
3 for sentence in tqdm(test_df['sentence2_preprocessed'].values):
4     vector = np.zeros(300)
5     cnt_words = 0
6     for word in sentence.split():
7         vector += ft.get_word_vector(word)
8         cnt_words += 1
9     if cnt_words != 0:
10        vector /= cnt_words
11    x_test_s2_avg_fasttext.append(vector)
12
13 print(len(x_test_s2_avg_fasttext))
14 print(len(x_test_s2_avg_fasttext[0]))

```

100% |██████████| 9824/9824 [00:00<00:00, 13819.69it/s]

9824

300

saving the fasttext ecoded array

In [29]:

```

1 x_train_s1_avg_fasttext = np.vstack(x_train_s1_avg_fasttext)
2 x_train_s2_avg_fasttext = np.vstack(x_train_s2_avg_fasttext)
3 x_val_s1_avg_fasttext = np.vstack(x_val_s1_avg_fasttext)
4 x_val_s2_avg_fasttext = np.vstack(x_val_s2_avg_fasttext)
5 x_test_s1_avg_fasttext = np.vstack(x_test_s1_avg_fasttext)
6 x_test_s2_avg_fasttext = np.vstack(x_test_s2_avg_fasttext)

```

In [30]:

```

1 np.save('x_train_s1_avg_fasttext.npy', x_train_s1_avg_fasttext)
2 np.save('x_train_s2_avg_fasttext.npy', x_train_s2_avg_fasttext)
3 np.save('x_val_s1_avg_fasttext.npy', x_val_s1_avg_fasttext)
4 np.save('x_val_s2_avg_fasttext.npy', x_val_s2_avg_fasttext)
5 np.save('x_test_s1_avg_fasttext.npy', x_test_s1_avg_fasttext)
6 np.save('x_test_s2_avg_fasttext.npy', x_test_s2_avg_fasttext)

```

In [31]:

```

1 x_train_s1_avg_fasttext = np.load('x_train_s1_avg_fasttext.npy', mmap_mode='r')
2 x_train_s2_avg_fasttext = np.load('x_train_s2_avg_fasttext.npy', mmap_mode='r')
3 x_val_s1_avg_fasttext = np.load('x_val_s1_avg_fasttext.npy', mmap_mode='r')
4 x_val_s2_avg_fasttext = np.load('x_val_s2_avg_fasttext.npy', mmap_mode='r')
5 x_test_s1_avg_fasttext = np.load('x_test_s1_avg_fasttext.npy', mmap_mode='r')
6 x_test_s2_avg_fasttext = np.load('x_test_s2_avg_fasttext.npy', mmap_mode='r')

```

Xgboost with fasttext embeddings

```
In [33]: 1 X_train = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val = val_df[to_keep]
4 y_val = val_df['gold_label']
5 x_train = np.hstack((X_train, x_train_s1_avg_fasttext, x_train_s2_avg_fasttext))
6 x_val = np.hstack((X_val, x_val_s1_avg_fasttext, x_val_s2_avg_fasttext))
```

```
In [34]: 1 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , va
2 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , valu
```

```
In [40]: 1 import xgboost as xgb
2 xgbc = xgb.XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=
3 colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
4 gamma=0.1, gpu_id=-1, importance_type=None,
5 interaction_constraints='', learning_rate=0.1, max_delta_step=
6 max_depth=12, min_child_weight=1, missing=np.nan,
7 monotone_constraints='()', n_estimators=200, n_jobs=16,
8 num_parallel_tree=1, objective='multi:softprob', predictor='au
9 random_state=0, reg_alpha=0, reg_lambda=100,
10 scale_pos_weight=None, subsample=1, tree_method='exact',
11 validate_parameters=1, verbosity=None, use_label_encoder=False)
```

```
In [41]: 1 xgbc.fit(x_train, y_train)
```

[11:36:17] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
Out[41]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
gamma=0.1, gpu_id=-1, importance_type=None,
interaction_constraints='', learning_rate=0.1, max_delta_step=0,
max_depth=12, min_child_weight=1, missing=nan,
monotone_constraints='()', n_estimators=200, n_jobs=16,
num_parallel_tree=1, objective='multi:softprob', predictor='aut
o',
random_state=0, reg_alpha=0, reg_lambda=100,
scale_pos_weight=None, subsample=1, tree_method='exact',
use_label_encoder=False, validate_parameters=1, verbosity=None)
```

```
In [42]: 1 val_pred = xgbc.predict(x_val)
2 from sklearn.metrics import accuracy_score
3 acc = accuracy_score(y_val, val_pred)
4 acc
```

```
Out[42]: 0.7050396260922577
```

we are getting 0.7 accuracy with average fasttext embeddings. No improvement. Let us try tfidf weighted fasttext

tfidf weighted fasttext embeddings

sentence1

```
In [72]: 1 #creating tfidf vecotrizer for sentence1  
2 from sklearn.feature_extraction.text import TfidfVectorizer  
3 tfidf_vectorizer_s1 = TfidfVectorizer(ngram_range=(1, 1))
```

```
In [73]: 1 #fitting on train data  
2 tfidf_vectorizer_s1.fit(train_df['sentence1_preprocessed'])
```

```
Out[73]: TfidfVectorizer()
```

```
In [74]: 1 #transforming train-test-val data  
2 train_s1_tfidf = tfidf_vectorizer_s1.transform(train_df['sentence1_preproces  
3 val_s1_tfidf = tfidf_vectorizer_s1.transform(val_df['sentence1_preprocessed'  
4 test_s1_tfidf = tfidf_vectorizer_s1.transform(test_df['sentence1_preprocesse
```

```
In [75]: 1 len(tfidf_vectorizer_s1.get_feature_names())
```

```
Out[75]: 17960
```

```
In [82]: 1 # creating a dic with keys as tokens and values and idf values  
2 idf_dict_s1 = dict(zip(tfidf_vectorizer_s1.get_feature_names(), tfidf_vector
```

```
In [138]: 1 def tfidf_weighted_fasttext(sentences, ft, idf_dict):
2
3     """
4         Given a list of sentences, the function returns
5         a 2d array where each row is tf-idf weighted fasttext
6         vector of each sentence.
7     args:
8         sentences: list of sentences
9         ft: a fasttext model
10        idf_dict: a dictionary with tokens as keys and idf values and values
11
12    """
13
14    from collections import Counter
15    sentences_tfidf_fasttext = list()
16    vocab = list(idf_dict.keys())
17    for sent in sentences:
18        words = sent.split(' ')
19        tf_dict = dict(Counter(words))
20        tf_idf_sum = 0
21        doc_vector = np.zeros(300)
22        for word in words:
23            if word in vocab:
24                fast_text = ft.get_word_vector(word)
25                idf = idf_dict[word]
26                tf = tf_dict[word]
27                tf_idf = tf*idf
28                weighted_word_vec = fast_text*tf_idf
29                tf_idf_sum = tf_idf_sum + tf_idf
30                doc_vector = doc_vector + weighted_word_vec
31
32            if tf_idf_sum != 0:
33                doc_vector = doc_vector/tf_idf_sum
34
35        sentences_tfidf_fasttext.append(doc_vector)
36
37    return np.vstack(sentences_tfidf_fasttext)
38
39
```

```
In [139]: 1 sentences = train_df['sentence1_preprocessed'][0:5].values
2 v = TfidfVectorizer(ngram_range=(1,1))
3 v.fit(sentences)
4 idf_dict = dict(zip(v.get_feature_names(), v.idf_))
```

```
In [140]: 1 result = tfidf_weighted_fasttext(sentences, ft, idf_dict)
```

```
In [141]: 1 result.shape
```

Out[141]: (5, 300)

Type *Markdown* and *LaTeX*: α^2

train sentence1 tfidf weighted fasttext

Type *Markdown* and *LaTeX*: α^2

```
In [142]: 1 sentences = train_df['sentence1_preprocessed'].values
           2 train_s1_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s1)
```

```
In [143]: 1 train_s1_tfidf_fasttext.shape
```

```
Out[143]: (549361, 300)
```

validation sentence1 tfidf weighted fasttext

```
In [144]: 1 sentences = val_df['sentence1_preprocessed'].values
           2 val_s1_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s1)
```

```
In [145]: 1 val_s1_tfidf_fasttext.shape
```

```
Out[145]: (9842, 300)
```

test sentence1 tfidf weighted fasttext

```
In [146]: 1 sentences = test_df['sentence1_preprocessed'].values
           2 test_s1_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s1)
```

```
In [147]: 1 test_s1_tfidf_fasttext.shape
```

```
Out[147]: (9824, 300)
```

sentence2

```
In [148]: 1 #creating tfidf vecotrizer for sentence2
           2 from sklearn.feature_extraction.text import TfidfVectorizer
           3 tfidf_vectorizer_s2 = TfidfVectorizer(ngram_range=(1, 1))
```

```
In [149]: 1 #fitting on train data
           2 tfidf_vectorizer_s2.fit(train_df['sentence2_preprocessed'])
```

```
Out[149]: TfidfVectorizer()
```

```
In [150]: 1 len(tfidf_vectorizer_s2.get_feature_names())
```

```
Out[150]: 30038
```

```
In [151]: 1 # creating a dic with keys as tokens and values and idf values
           2 idf_dict_s2 = dict(zip(tfidf_vectorizer_s2.get_feature_names(), tfidf_vector
```

train sentence2 tfidf weighted fasttext

```
In [153]: 1 sentences = train_df['sentence2_preprocessed'].values
           2 train_s2_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s2)

In [154]: 1 train_s2_tfidf_fasttext.shape

Out[154]: (549361, 300)
```

validation sentence2 tfidf weighted fasttext

```
In [155]: 1 sentences = val_df['sentence2_preprocessed'].values
           2 val_s2_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s2)

In [157]: 1 val_s2_tfidf_fasttext.shape

Out[157]: (9842, 300)
```

test sentence2 tfidf weighted fasttext

```
In [158]: 1 sentences = test_df['sentence2_preprocessed'].values
           2 test_s2_tfidf_fasttext = tfidf_weighted_fasttext(sentences, ft, idf_dict_s2)

In [159]: 1 test_s2_tfidf_fasttext.shape

Out[159]: (9824, 300)

In [160]: 1 #saving the matrices
           2 np.save('train_s1_tfidf_fasttext.npy', train_s1_tfidf_fasttext)
           3 np.save('train_s2_tfidf_fasttext.npy', train_s2_tfidf_fasttext)
           4 np.save('val_s1_tfidf_fasttext.npy', val_s1_tfidf_fasttext)
           5 np.save('val_s2_tfidf_fasttext.npy', val_s2_tfidf_fasttext)
           6 np.save('test_s1_tfidf_fasttext.npy', test_s1_tfidf_fasttext)
           7 np.save('test_s2_tfidf_fasttext.npy', test_s2_tfidf_fasttext)

In [161]: 1 # Loading
           2 train_s1_tfidf_fasttext = np.load('train_s1_tfidf_fasttext.npy', mmap_mode='r')
           3 train_s2_tfidf_fasttext = np.load('train_s2_tfidf_fasttext.npy', mmap_mode='r')
           4 val_s1_tfidf_fasttext = np.load('val_s1_tfidf_fasttext.npy', mmap_mode='r')
           5 val_s2_tfidf_fasttext = np.load('val_s2_tfidf_fasttext.npy', mmap_mode='r')
           6 test_s1_tfidf_fasttext = np.load('test_s1_tfidf_fasttext.npy', mmap_mode='r')
           7 test_s2_tfidf_fasttext = np.load('test_s2_tfidf_fasttext.npy', mmap_mode='r')
```

```
In [164]: 1 X_train = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val = val_df[to_keep]
4 y_val = val_df['gold_label']
5 x_train = np.hstack((X_train, train_s1_tfidf_fasttext, train_s2_tfidf_fasttext))
6 x_val = np.hstack((X_val, val_s1_tfidf_fasttext, val_s2_tfidf_fasttext))
```

```
In [ ]: 1 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , va
2 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , valu
```

```
In [165]: 1 import xgboost as xgb
2 xgbc = xgb.XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=
3 colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
4 gamma=0.1, gpu_id=-1, importance_type=None,
5 interaction_constraints='', learning_rate=0.1, max_delta_step=
6 max_depth=12, min_child_weight=1, missing=np.nan,
7 monotone_constraints='()', n_estimators=200, n_jobs=16,
8 num_parallel_tree=1, objective='multi:softprob', predictor='au
9 random_state=0, reg_alpha=0, reg_lambda=100,
10 scale_pos_weight=None, subsample=1, tree_method='exact',
11 validate_parameters=1, verbosity=None, use_label_encoder=False)
```

```
In [166]: 1 xgbc.fit(x_train, y_train)
```

[18:22:41] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
Out[166]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
gamma=0.1, gpu_id=-1, importance_type=None,
interaction_constraints='', learning_rate=0.1, max_delta_step=0,
max_depth=12, min_child_weight=1, missing=nan,
monotone_constraints='()', n_estimators=200, n_jobs=16,
num_parallel_tree=1, objective='multi:softprob', predictor='aut
o',
random_state=0, reg_alpha=0, reg_lambda=100,
scale_pos_weight=None, subsample=1, tree_method='exact',
use_label_encoder=False, validate_parameters=1, verbosity=None)
```

```
In [167]: 1 val_pred = xgbc.predict(x_val)
2 from sklearn.metrics import accuracy_score
3 acc = accuracy_score(y_val, val_pred)
4 acc
```

```
Out[167]: 0.7159114001219264
```

little improvement after taking tfidf weighted fasttext embeddings

In [4]:

```

1 from prettytable import PrettyTable
2 myTable = PrettyTable(["Feature set", "Model", "accuracy"])
3 myTable.add_row(["Only extracted features", "RandomForest", "0.58"])
4 myTable.add_row(["Extracted features + avg GloVe embeddings", "RandomForest", "0.66"])
5 myTable.add_row(["Extracted features + avg GloVe embeddings", "XGBoost", "0.7"])
6 myTable.add_row(["Extracted features + avg fasttext embeddings", "XGBoost", "0.7"])
7 myTable.add_row(["Extracted features + tfidf weighted fasttext embeddings", "XGBoost", "0.71"])
8 print('Summary of classical machine learning aproaches:')
9 print(myTable)

```

Summary of classical machine learning aproaches:

	Feature set	Model	accuracy
58	Only extracted features	RandomForest	0.58
66	Extracted features + avg GloVe embeddings	RandomForest	0.66
70	Extracted features + avg GloVe embeddings	XGBoost	0.7
70	Extracted features + avg fasttext embeddings	XGBoost	0.7
71	Extracted features + tfidf weighted fasttext embeddings	XGBoost	0.71

Moving towards Deep Learning

In [3]:

```

1 from tensorflow.keras.layers import Input, Dense, concatenate, BatchNormalization
2 from tensorflow.keras.models import Model
3 from tensorflow.keras.utils import plot_model
4 from tensorflow.keras.utils import to_categorical
5 from tensorflow.keras.regularizers import L1, L2
6 from tensorflow.keras.optimizers import Adam, SGD, RMSprop

```

In [4]:

```

1 import tensorflow as tf
2 tf.config.experimental.allow_growth = True
3

```

In [5]:

```

1 # Loading the dataframes in which features have already been extracted
2 train_df = pd.read_csv("train_df.csv")
3 val_df = pd.read_csv("val_df.csv")
4 test_df = pd.read_csv("test_df.csv")

```

```
In [6]: 1 # Loading avg glove vectors
2 x_train_s1_tfidf2v = np.load('x_train_s1_tfidf2v.npy', mmap_mode='r')
3 x_train_s2_tfidf2v = np.load('x_train_s2_tfidf2v.npy', mmap_mode='r')
4 x_val_s1_tfidf2v = np.load('x_val_s1_tfidf2v.npy', mmap_mode='r')
5 x_val_s2_tfidf2v = np.load('x_val_s2_tfidf2v.npy', mmap_mode='r')
6 # x_test_s1_tfidf2v = np.load('x_test_s1_tfidf2v.npy', mmap_mode='r')
7 # x_test_s2_tfidf2v = np.load('x_test_s2_tfidf2v.npy', mmap_mode='r')

In [6]: 1 # List of columns to be kept from loaded dataframes
2 to_keep = ['log_bleu_score', 'sum_polarity', 'diff_subjectivity',
3             'cosine_similarity', 'tfidf_cosine_sim', 'diff_num_words', 's1_CC',
4             's1_IN', 's1_PRP', 's1_PRP$', 's1_RB', 's1_RBR', 's1_RBS', 's1_UH', 's
5             's2_PRP$', 's2_RB', 's2_RBR', 's2_RBS', 's2_UH', 'ratio_noun', 'ratio
6             'diff_adjective', 'levenshtein_dist', 'tfidf_w2v_cosine',
7             'tfidf_cosine_sim_ngram']

In [7]: 1 #selecting required columns
2 X_train_1 = train_df[to_keep]
3 y_train = train_df['gold_label']
4 X_val_1 = val_df[to_keep]
5 y_val = val_df['gold_label']

In [8]: 1 #converting target categories to numbers: neutral:0, contradiction:1, entailment:2
2 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , value=0)
3 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , value=0)

In [9]: 1 #one hot encoding target
2 y_train_enc = to_categorical(y_train)
3 y_val_enc = to_categorical(y_val)
```

model0: considering only avg glove embeddings

```
In [12]: 1 input2 = Input(shape=(300,)) #for sentence1 embeddings
2 input3 = Input(shape=(300,)) #for sentence2 embeddings
3 merged = concatenate([input2, input3]) #merging all inputs
4 norm1 = BatchNormalization()(merged)
5 dense1 = Dense(units = 50, activation = 'relu')(norm1)
6 dense2 = Dense(units = 40, activation = 'relu')(dense1)
7 dense3 = Dense(units = 30, activation = 'relu')(dense2)
8 norm2 = BatchNormalization()(dense3)
9 dense4 = Dense(units = 20, activation = 'relu')(norm2)
10 dense5 = Dense(units = 10, activation = 'relu')(dense4)
11 output = Dense(units = 3, activation = 'softmax')(dense5)
12 model0 = Model(inputs = [input2, input3], outputs = output)

In [13]: 1 model0.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric
```

In [14]:

```
1 model0.fit([x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_enc, epochs =  
2 validation_data = ([x_val_s1_tfidf2v, x_val_s2_tfidf2v], y_val_
```

Epoch 1/50
4292/4292 [=====] - 20s 4ms/step - loss: 0.8394 - accuracy: 0.6137 - val_loss: 0.7503 - val_accuracy: 0.6715
Epoch 2/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.7406 - accuracy: 0.6768 - val_loss: 0.7123 - val_accuracy: 0.6944
Epoch 3/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.7113 - accuracy: 0.6929 - val_loss: 0.6980 - val_accuracy: 0.7017
Epoch 4/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.6938 - accuracy: 0.7021 - val_loss: 0.6847 - val_accuracy: 0.7038
Epoch 5/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.6816 - accuracy: 0.7092 - val_loss: 0.6860 - val_accuracy: 0.7074
Epoch 6/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.6731 - accuracy: 0.7139 - val_loss: 0.6701 - val_accuracy: 0.7166
Epoch 7/50
4292/4292 [=====] - 18s 4ms/step - loss: 0.6654 - accuracy: 0.7182 - val_loss: 0.6664 - val_accuracy: 0.7182

```
In [15]: 1 model0.fit([x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_enc, epochs =
2           validation_data = ([x_val_s1_tfidf2v, x_val_s2_tfidf2v], y_val_)

Epoch 1/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6099 - accuracy: 0.7463 - val_loss: 0.6453 - val_accuracy: 0.7322
Epoch 2/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6092 - accuracy: 0.7466 - val_loss: 0.6461 - val_accuracy: 0.7303
Epoch 3/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6093 - accuracy: 0.7460 - val_loss: 0.6455 - val_accuracy: 0.7280
Epoch 4/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6088 - accuracy: 0.7466 - val_loss: 0.6423 - val_accuracy: 0.7314
Epoch 5/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6084 - accuracy: 0.7471 - val_loss: 0.6435 - val_accuracy: 0.7311
Epoch 6/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6082 - accuracy: 0.7470 - val_loss: 0.6428 - val_accuracy: 0.7304
Epoch 7/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6078 - accuracy: 0.7476 - val_loss: 0.6431 - val_accuracy: 0.7306
Epoch 8/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6075 - accuracy: 0.7472 - val_loss: 0.6444 - val_accuracy: 0.7309
Epoch 9/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6067 - accuracy: 0.7472 - val_loss: 0.6431 - val_accuracy: 0.7307
Epoch 10/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6063 - accuracy: 0.7478 - val_loss: 0.6485 - val_accuracy: 0.7321
Epoch 11/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6063 - accuracy: 0.7479 - val_loss: 0.6466 - val_accuracy: 0.7306
Epoch 12/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6058 - accuracy: 0.7485 - val_loss: 0.6424 - val_accuracy: 0.7342
Epoch 13/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6058 - accuracy: 0.7484 - val_loss: 0.6394 - val_accuracy: 0.7344
Epoch 14/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6060 - accuracy: 0.7478 - val_loss: 0.6458 - val_accuracy: 0.7332
Epoch 15/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6053 - accuracy: 0.7480 - val_loss: 0.6461 - val_accuracy: 0.7303
Epoch 16/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6045 - accuracy: 0.7482 - val_loss: 0.6459 - val_accuracy: 0.7296
Epoch 17/20
4292/4292 [=====] - 19s 4ms/step - loss: 0.6047 - accuracy: 0.7484 - val_loss: 0.6429 - val_accuracy: 0.7326
Epoch 18/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6044 - accuracy: 0.7485 - val_loss: 0.6416 - val_accuracy: 0.7322
```

```
Epoch 19/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6040 - accuracy: 0.7491 - val_loss: 0.6417 - val_accuracy: 0.7320
Epoch 20/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.6041 - accuracy: 0.7491 - val_loss: 0.6388 - val_accuracy: 0.7357
```

Out[15]: <keras.callbacks.History at 0x213f4fdd780>

```
In [16]: 1 model0.fit([x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_enc, epochs =
2           validation_data = ([x_val_s1_tfidf2v, x_val_s2_tfidf2v], y_val_)

Epoch 1/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6038 - accuracy: 0.7489 - val_loss: 0.6420 - val_accuracy: 0.7335
Epoch 2/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6033 - accuracy: 0.7493 - val_loss: 0.6446 - val_accuracy: 0.7331
Epoch 3/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6033 - accuracy: 0.7487 - val_loss: 0.6431 - val_accuracy: 0.7311
Epoch 4/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6030 - accuracy: 0.7497 - val_loss: 0.6425 - val_accuracy: 0.7337
Epoch 5/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6028 - accuracy: 0.7500 - val_loss: 0.6430 - val_accuracy: 0.7315
Epoch 6/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6025 - accuracy: 0.7494 - val_loss: 0.6432 - val_accuracy: 0.7283
Epoch 7/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6023 - accuracy: 0.7499 - val_loss: 0.6439 - val_accuracy: 0.7282
Epoch 8/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6021 - accuracy: 0.7493 - val_loss: 0.6427 - val_accuracy: 0.7312
Epoch 9/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6015 - accuracy: 0.7500 - val_loss: 0.6424 - val_accuracy: 0.7322
Epoch 10/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6013 - accuracy: 0.7507 - val_loss: 0.6465 - val_accuracy: 0.7327
Epoch 11/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6012 - accuracy: 0.7502 - val_loss: 0.6449 - val_accuracy: 0.7331
Epoch 12/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6008 - accuracy: 0.7506 - val_loss: 0.6412 - val_accuracy: 0.7324
Epoch 13/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6008 - accuracy: 0.7503 - val_loss: 0.6396 - val_accuracy: 0.7313
Epoch 14/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6012 - accuracy: 0.7501 - val_loss: 0.6441 - val_accuracy: 0.7315
Epoch 15/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6006 - accuracy: 0.7507 - val_loss: 0.6449 - val_accuracy: 0.7316
Epoch 16/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.5999 - accuracy: 0.7500 - val_loss: 0.6460 - val_accuracy: 0.7312
Epoch 17/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6002 - accuracy: 0.7505 - val_loss: 0.6446 - val_accuracy: 0.7342
Epoch 18/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.6001 - accuracy: 0.7502 - val_loss: 0.6418 - val_accuracy: 0.7337
```

```

Epoch 19/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.5997 - accuracy: 0.7510 - val_loss: 0.6398 - val_accuracy: 0.7308
Epoch 20/20
4292/4292 [=====] - 18s 4ms/step - loss: 0.5997 - accuracy: 0.7509 - val_loss: 0.6390 - val_accuracy: 0.7351

```

Out[16]: <keras.callbacks.History at 0x213f4fd0ac8>

Model0_a

In [31]:	<pre> 1 input2 = Input(shape=(300,)) #for sentence1 embeddings 2 norm2 = BatchNormalization()(input2) 3 input3 = Input(shape=(300,)) #for sentence2 embeddings 4 norm3 = BatchNormalization()(input3) 5 merged = concatenate([norm2, norm3]) #merging all inputs #norm1 = BatchNormalization()(merged) 6 dense1 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0 7 drop1 = Dropout(0.01)(dense1) 8 dense2 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0 9 drop2 = Dropout(0.01)(dense2) 10 dense3 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0 11 dense4 = Dense(units = 50, activation = 'relu')(dense3) 12 output = Dense(units = 3, activation = 'softmax')(dense4) 13 model0_a = Model(inputs = [input2, input3], outputs = output) </pre>
----------	--

In [32]:	<pre>1 model0_a.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metr</pre>
----------	---

In [34]:

```

1 tf.debugging.set_log_device_placement(True)
2
3 # Place tensors on the CPU
4 with tf.device('/CPU:0'):
5     model0_a.fit([x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_enc, ep_
6                 validation_data = ([x_val_s1_tfidf2v, x_val_s2_tfidf2v], y_val_

```

```

Epoch 1/100
1073/1073 [=====] - 11s 8ms/step - loss: 0.8531 - accuracy: 0.6338 - val_loss: 0.7762 - val_accuracy: 0.6831
Epoch 2/100
1073/1073 [=====] - 8s 7ms/step - loss: 0.7543 - accuracy: 0.6929 - val_loss: 0.7307 - val_accuracy: 0.7064
Epoch 3/100
1073/1073 [=====] - 8s 7ms/step - loss: 0.7202 - accuracy: 0.7111 - val_loss: 0.7142 - val_accuracy: 0.7123
Epoch 4/100
1073/1073 [=====] - 8s 7ms/step - loss: 0.6995 - accuracy: 0.7226 - val_loss: 0.6951 - val_accuracy: 0.7253
Epoch 5/100
1073/1073 [=====] - 8s 7ms/step - loss: 0.6845 - accuracy: 0.7308 - val_loss: 0.6869 - val_accuracy: 0.7301
Epoch 6/100
1073/1073 [=====] - 8s 8ms/step - loss: 0.6736 - accuracy: 0.7368 - val_loss: 0.6803 - val_accuracy: 0.7350
Epoch 7/100
1073/1073 [=====] - 8s 8ms/step - loss: 0.6636 - accuracy: 0.7411 - val_loss: 0.6754 - val_accuracy: 0.7399

```

model1: considering extracted features+avg glove embeddings

In [16]:

```

1 input1 = Input(shape=(29,)) #for extracted features
2 input2 = Input(shape=(300,)) #for sentence1 embeddings
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 merged = concatenate([input1, input2, input3]) #merging all inputs
5 norm1 = BatchNormalization()(merged)
6 dense1 = Dense(units = 50, activation = 'relu')(norm1)
7 dense2 = Dense(units = 40, activation = 'relu')(dense1)
8 dense3 = Dense(units = 30, activation = 'relu')(dense2)
9 norm2 = BatchNormalization()(dense3)
10 dense4 = Dense(units = 20, activation = 'relu')(norm2)
11 dense5 = Dense(units = 10, activation = 'relu')(dense4)
12 output = Dense(units = 3, activation = 'softmax')(dense5)
13 model1 = Model(inputs = [input1, input2, input3], outputs = output)

```

In [17]:

```
1 model1.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric
```

```
In [18]: 1 model1.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_en  
2           validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v])
```

```
Epoch 1/50  
4292/4292 [=====] - 23s 5ms/step - loss: 0.7971 - accuracy: 0.6417 - val_loss: 0.7009 - val_accuracy: 0.6971  
Epoch 2/50  
4292/4292 [=====] - 17s 4ms/step - loss: 0.7027 - accuracy: 0.6977 - val_loss: 0.6650 - val_accuracy: 0.7143  
Epoch 3/50  
4292/4292 [=====] - 17s 4ms/step - loss: 0.6749 - accuracy: 0.7125 - val_loss: 0.6568 - val_accuracy: 0.7214  
Epoch 4/50  
4292/4292 [=====] - 18s 4ms/step - loss: 0.6584 - accuracy: 0.7211 - val_loss: 0.6435 - val_accuracy: 0.7313  
Epoch 5/50  
4292/4292 [=====] - 17s 4ms/step - loss: 0.6469 - accuracy: 0.7266 - val_loss: 0.6365 - val_accuracy: 0.7297  
Epoch 6/50  
4292/4292 [=====] - 17s 4ms/step - loss: 0.6388 - accuracy: 0.7309 - val_loss: 0.6262 - val_accuracy: 0.7379  
Epoch 7/50  
4292/4292 [=====] - 17s 4ms/step - loss: 0.6322 - accuracy: 0.7333 - val_loss: 0.6202 - val_accuracy: 0.7402
```

In [19]:

```
1 model1.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_en  
2 validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v]
```

```
Epoch 1/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5752 - accuracy: 0.7624 - val_loss: 0.5992 - val_accuracy: 0.7520  
Epoch 2/20  
4292/4292 [=====] - 19s 4ms/step - loss: 0.5748 - accuracy: 0.7624 - val_loss: 0.5974 - val_accuracy: 0.7552  
Epoch 3/20  
4292/4292 [=====] - 18s 4ms/step - loss: 0.5749 - accuracy: 0.7620 - val_loss: 0.5993 - val_accuracy: 0.7522  
Epoch 4/20  
4292/4292 [=====] - 19s 4ms/step - loss: 0.5744 - accuracy: 0.7625 - val_loss: 0.5978 - val_accuracy: 0.7555  
Epoch 5/20  
4292/4292 [=====] - 19s 4ms/step - loss: 0.5738 - accuracy: 0.7626 - val_loss: 0.5985 - val_accuracy: 0.7526  
Epoch 6/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5733 - accuracy: 0.7629 - val_loss: 0.6009 - val_accuracy: 0.7555  
Epoch 7/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5731 - accuracy: 0.7636 - val_loss: 0.6005 - val_accuracy: 0.7534  
Epoch 8/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5728 - accuracy: 0.7630 - val_loss: 0.6010 - val_accuracy: 0.7510  
Epoch 9/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5721 - accuracy: 0.7635 - val_loss: 0.5982 - val_accuracy: 0.7551  
Epoch 10/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5719 - accuracy: 0.7642 - val_loss: 0.6077 - val_accuracy: 0.7534  
Epoch 11/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5711 - accuracy: 0.7644 - val_loss: 0.6005 - val_accuracy: 0.7516  
Epoch 12/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5715 - accuracy: 0.7639 - val_loss: 0.5984 - val_accuracy: 0.7543  
Epoch 13/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5708 - accuracy: 0.7646 - val_loss: 0.5962 - val_accuracy: 0.7537  
Epoch 14/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5712 - accuracy: 0.7648 - val_loss: 0.6000 - val_accuracy: 0.7539  
Epoch 15/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5702 - accuracy: 0.7650 - val_loss: 0.6015 - val_accuracy: 0.7525  
Epoch 16/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5701 - accuracy: 0.7648 - val_loss: 0.6028 - val_accuracy: 0.7577  
Epoch 17/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5699 - accuracy: 0.7645 - val_loss: 0.5993 - val_accuracy: 0.7561  
Epoch 18/20  
4292/4292 [=====] - 17s 4ms/step - loss: 0.5694 - accuracy: 0.7647 - val_loss: 0.5994 - val_accuracy: 0.7552
```

```

Epoch 19/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.5693 - accuracy: 0.7647 - val_loss: 0.5987 - val_accuracy: 0.7546
Epoch 20/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.5693 - accuracy: 0.7645 - val_loss: 0.5980 - val_accuracy: 0.7571

```

Out[19]: <keras.callbacks.History at 0x2aca478d860>

model1_a

```

In [61]: 1 input1 = Input(shape=(29,)) #for extracted features
2 input2 = Input(shape=(300,)) #for sentence1 embeddings
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 merged = concatenate([input1, input2, input3]) #merging all inputs
5 norm1 = BatchNormalization()(merged)
6 dense1 = Dense(units = 50, activation = 'relu')(merged)
7 dense2 = Dense(units = 50, activation = 'relu')(dense1)
8 dense3 = Dense(units = 50, activation = 'relu')(dense2)
9 output = Dense(units = 3, activation = 'softmax')(dense3)
10 model1_a = Model(inputs = [input1, input2, input3], outputs = output)

```

```

In [62]: 1 model1_a.compile(optimizer = Adam(0.001), loss = 'categorical_crossentropy',
2 tf.debugging.set_log_device_placement(True)
3
4 # Place tensors on the CPU
5 with tf.device('/CPU:0'):
6     model1_a.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train,
7                  validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v]

```

```

Epoch 1/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.9425 - accuracy: 0.5589 - val_loss: 0.8820 - val_accuracy: 0.5935
Epoch 2/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.8529 - accuracy: 0.6141 - val_loss: 0.8227 - val_accuracy: 0.6304
Epoch 3/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.8321 - accuracy: 0.6260 - val_loss: 0.8034 - val_accuracy: 0.6467
Epoch 4/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.8176 - accuracy: 0.6335 - val_loss: 0.7844 - val_accuracy: 0.6520
Epoch 5/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.8003 - accuracy: 0.6437 - val_loss: 0.8355 - val_accuracy: 0.6159
Epoch 6/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.7881 - accuracy: 0.6508 - val_loss: 0.7705 - val_accuracy: 0.6612
Epoch 7/100
1073/1073 [=====] - 4s 4ms/step - loss: 0.7753 - accuracy: 0.6571 - val_loss: 0.7558 - val_accuracy: 0.6712

```

```
In [63]: 1 with tf.device('/CPU:0'):
2     model1_a.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_tr
3                 validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v

Epoch 1/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6231 - accuracy: 0.7374 - val_loss: 0.6644 - val_accuracy: 0.7195
Epoch 2/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6219 - accuracy: 0.7382 - val_loss: 0.6702 - val_accuracy: 0.7167
Epoch 3/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6216 - accuracy: 0.7384 - val_loss: 0.6611 - val_accuracy: 0.7228
Epoch 4/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6206 - accuracy: 0.7394 - val_loss: 0.6817 - val_accuracy: 0.7150
Epoch 5/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6218 - accuracy: 0.7385 - val_loss: 0.6624 - val_accuracy: 0.7214
Epoch 6/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6214 - accuracy: 0.7385 - val_loss: 0.6626 - val_accuracy: 0.7198
Epoch 7/50
1073/1073 [=====] - 4s 4ms/step - loss: 0.6211 - accuracy: 0.7385 - val_loss: 0.6626 - val_accuracy: 0.7198
```

model1_b

```
In [11]: 1 input1 = Input(shape=(29,)) #for extracted features
2 input2 = Input(shape=(300,)) #for sentence1 embeddings
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 merged = concatenate([input1, input2, input3]) #merging all inputs
5 norm1 = BatchNormalization()(merged)
6 dense1 = Dense(units = 70, activation = 'relu')(norm1)
7 dense2 = Dense(units = 60, activation = 'relu')(dense1)
8 dense3 = Dense(units = 50, activation = 'relu')(dense2)
9 dense4 = Dense(units = 40, activation = 'relu')(dense3)
10 dense5 = Dense(units = 30, activation = 'relu')(dense4)
11 dense6 = Dense(units = 20, activation = 'relu')(dense5)
12 dense7 = Dense(units = 10, activation = 'relu')(dense6)
13 output = Dense(units = 3, activation = 'softmax')(dense7)
14 model1_b = Model(inputs = [input1, input2, input3], outputs = output)
```

```
In [12]: 1 model1_b.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metr  
2 model1_b.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_  
3 validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v])  
  
Epoch 1/100  
4292/4292 [=====] - 31s 7ms/step - loss: 0.7687 - ac  
curacy: 0.6592 - val_loss: 0.6857 - val_accuracy: 0.7053  
Epoch 2/100  
4292/4292 [=====] - 25s 6ms/step - loss: 0.6825 - ac  
curacy: 0.7081 - val_loss: 0.6555 - val_accuracy: 0.7228  
Epoch 3/100  
4292/4292 [=====] - 22s 5ms/step - loss: 0.6551 - ac  
curacy: 0.7228 - val_loss: 0.6410 - val_accuracy: 0.7342  
Epoch 4/100  
4292/4292 [=====] - 21s 5ms/step - loss: 0.6375 - ac  
curacy: 0.7314 - val_loss: 0.6333 - val_accuracy: 0.7318  
Epoch 5/100  
4292/4292 [=====] - 23s 5ms/step - loss: 0.6251 - ac  
curacy: 0.7381 - val_loss: 0.6214 - val_accuracy: 0.7413  
Epoch 6/100  
4292/4292 [=====] - 23s 5ms/step - loss: 0.6158 - ac  
curacy: 0.7428 - val_loss: 0.6189 - val_accuracy: 0.7393  
Epoch 7/100  
4292/4292 [=====] - 23s 5ms/step - loss: 0.6158 - ac  
curacy: 0.7428 - val_loss: 0.6189 - val_accuracy: 0.7393
```

```
In [13]: 1 from keras import backend as K  
2 K.set_value(model1_b.optimizer.learning_rate, 0.0001)
```

In [14]:

```
1 model1_b.fit([X_train_1, x_train_s1_tfidf2v, x_train_s2_tfidf2v], y_train_
2 validation_data = ([X_val_1, x_val_s1_tfidf2v, x_val_s2_tfidf2v]
```

```
Epoch 1/20
4292/4292 [=====] - 17s 4ms/step - loss: 0.5067 - accuracy: 0.7943 - val_loss: 0.6106 - val_accuracy: 0.7558
Epoch 2/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.5030 - accuracy: 0.7957 - val_loss: 0.6165 - val_accuracy: 0.7540
Epoch 3/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.5018 - accuracy: 0.7963 - val_loss: 0.6167 - val_accuracy: 0.7544
Epoch 4/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.5012 - accuracy: 0.7962 - val_loss: 0.6164 - val_accuracy: 0.7552
Epoch 5/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.5002 - accuracy: 0.7971 - val_loss: 0.6217 - val_accuracy: 0.7525
Epoch 6/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.5000 - accuracy: 0.7970 - val_loss: 0.6203 - val_accuracy: 0.7544
Epoch 7/20
4292/4292 [=====] - 16s 4ms/step - loss: 0.4994 - accuracy: 0.7976 - val_loss: 0.6193 - val_accuracy: 0.7537
Epoch 8/20
4292/4292 [=====] - 16s 4ms/step - loss: 0.4998 - accuracy: 0.7973 - val_loss: 0.6187 - val_accuracy: 0.7537
Epoch 9/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4992 - accuracy: 0.7976 - val_loss: 0.6171 - val_accuracy: 0.7535
Epoch 10/20
4292/4292 [=====] - 16s 4ms/step - loss: 0.4991 - accuracy: 0.7974 - val_loss: 0.6229 - val_accuracy: 0.7535
Epoch 11/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4983 - accuracy: 0.7981 - val_loss: 0.6209 - val_accuracy: 0.7542
Epoch 12/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4983 - accuracy: 0.7979 - val_loss: 0.6199 - val_accuracy: 0.7543
Epoch 13/20
4292/4292 [=====] - 16s 4ms/step - loss: 0.4985 - accuracy: 0.7977 - val_loss: 0.6202 - val_accuracy: 0.7552
Epoch 14/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4984 - accuracy: 0.7979 - val_loss: 0.6233 - val_accuracy: 0.7540
Epoch 15/20
4292/4292 [=====] - 16s 4ms/step - loss: 0.4979 - accuracy: 0.7982 - val_loss: 0.6217 - val_accuracy: 0.7536
Epoch 16/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4979 - accuracy: 0.7979 - val_loss: 0.6231 - val_accuracy: 0.7541
Epoch 17/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4976 - accuracy: 0.7983 - val_loss: 0.6224 - val_accuracy: 0.7537
Epoch 18/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4977 - accuracy: 0.7982 - val_loss: 0.6253 - val_accuracy: 0.7544
```

```

Epoch 19/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4976 - accuracy: 0.7981 - val_loss: 0.6231 - val_accuracy: 0.7531
Epoch 20/20
4292/4292 [=====] - 15s 4ms/step - loss: 0.4978 - accuracy: 0.7981 - val_loss: 0.6241 - val_accuracy: 0.7549

```

Out[14]: <keras.callbacks.History at 0x1d0a5d7e978>

Model2: using only average fasttext embeddings

In [14]:

```

1 x_train_s1_avg_fasttext = np.load('x_train_s1_avg_fasttext.npy', mmap_mode='r')
2 x_train_s2_avg_fasttext = np.load('x_train_s2_avg_fasttext.npy', mmap_mode='r')
3 x_val_s1_avg_fasttext = np.load('x_val_s1_avg_fasttext.npy', mmap_mode='r')
4 x_val_s2_avg_fasttext = np.load('x_val_s2_avg_fasttext.npy', mmap_mode='r')
5 x_test_s1_avg_fasttext = np.load('x_test_s1_avg_fasttext.npy', mmap_mode='r')
6 x_test_s2_avg_fasttext = np.load('x_test_s2_avg_fasttext.npy', mmap_mode='r')

```

In [16]:

```

1 input2 = Input(shape=(300,)) #for sentence1 embeddings
2 norm2 = BatchNormalization()(input2)
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 norm3 = BatchNormalization()(input3)
5 merged = concatenate([norm2, norm3]) #merging all inputs
#norm1 = BatchNormalization()(merged)
6 dense1 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0
7 drop1 = Dropout(0.01)(dense1)
8 dense2 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0
9 drop2 = Dropout(0.01)(dense2)
10 dense3 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.0
11 dense4 = Dense(units = 50, activation = 'relu')(dense3)
12 output = Dense(units = 3, activation = 'softmax')(dense4)
13 model2_a = Model(inputs = [input2, input3], outputs = output)

```

In [18]:

```
1 model2_a.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metr
2
3 tf.debugging.set_log_device_placement(True)
4
5 # Place tensors on the CPU
6 with tf.device('/CPU:0'):
7     model2_a.fit([x_train_s1_avg_fasttext, x_train_s2_avg_fasttext], y_train,
8                 validation_data = ([x_val_s1_avg_fasttext, x_val_s2_avg_fasttext])
```

Epoch 1/100
4292/4292 [=====] - 23s 5ms/step - loss: 0.8644 - accuracy: 0.6295 - val_loss: 0.7856 - val_accuracy: 0.6819
Epoch 2/100
4292/4292 [=====] - 21s 5ms/step - loss: 0.7768 - accuracy: 0.6826 - val_loss: 0.7426 - val_accuracy: 0.7025
Epoch 3/100
4292/4292 [=====] - 20s 5ms/step - loss: 0.7464 - accuracy: 0.7010 - val_loss: 0.7265 - val_accuracy: 0.7112
Epoch 4/100
4292/4292 [=====] - 20s 5ms/step - loss: 0.7280 - accuracy: 0.7114 - val_loss: 0.7223 - val_accuracy: 0.7166
Epoch 5/100
4292/4292 [=====] - 20s 5ms/step - loss: 0.7162 - accuracy: 0.7192 - val_loss: 0.7038 - val_accuracy: 0.7273
Epoch 6/100
4292/4292 [=====] - 20s 5ms/step - loss: 0.7073 - accuracy: 0.7241 - val_loss: 0.7044 - val_accuracy: 0.7298
Epoch 7/100
4292/4292 [=====] - 20s 5ms/step - loss: 0.7011 - accuracy: 0.7241 - val_loss: 0.7011 - val_accuracy: 0.7298

In [19]:

```
1 from keras import backend as K
2 K.set_value(model2_a.optimizer.learning_rate, 0.0001)
```

In [20]:

```
1 with tf.device('/CPU:0'):
2     model2_a.fit([x_train_s1_avg_fasttext, x_train_s2_avg_fasttext], y_train,
3                  validation_data = ([x_val_s1_avg_fasttext, x_val_s2_avg_fasttext])
Epoch 1/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6116 - accuracy: 0.7845 - val_loss: 0.6761 - val_accuracy: 0.7562
Epoch 2/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6013 - accuracy: 0.7890 - val_loss: 0.6778 - val_accuracy: 0.7566
Epoch 3/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5973 - accuracy: 0.7901 - val_loss: 0.6782 - val_accuracy: 0.7568
Epoch 4/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5947 - accuracy: 0.7915 - val_loss: 0.6782 - val_accuracy: 0.7579
Epoch 5/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5934 - accuracy: 0.7912 - val_loss: 0.6788 - val_accuracy: 0.7576
Epoch 6/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5918 - accuracy: 0.7922 - val_loss: 0.6771 - val_accuracy: 0.7573
Epoch 7/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5903 - accuracy: 0.7930 - val_loss: 0.6789 - val_accuracy: 0.7590
Epoch 8/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5897 - accuracy: 0.7930 - val_loss: 0.6780 - val_accuracy: 0.7571
Epoch 9/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5884 - accuracy: 0.7937 - val_loss: 0.6797 - val_accuracy: 0.7588
Epoch 10/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5873 - accuracy: 0.7940 - val_loss: 0.6799 - val_accuracy: 0.7568
Epoch 11/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5869 - accuracy: 0.7941 - val_loss: 0.6811 - val_accuracy: 0.7574
Epoch 12/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5861 - accuracy: 0.7948 - val_loss: 0.6798 - val_accuracy: 0.7569
Epoch 13/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5851 - accuracy: 0.7951 - val_loss: 0.6813 - val_accuracy: 0.7559
Epoch 14/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5850 - accuracy: 0.7946 - val_loss: 0.6817 - val_accuracy: 0.7564
Epoch 15/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5842 - accuracy: 0.7948 - val_loss: 0.6841 - val_accuracy: 0.7552
Epoch 16/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5839 - accuracy: 0.7947 - val_loss: 0.6820 - val_accuracy: 0.7580
Epoch 17/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5830 - accuracy: 0.7954 - val_loss: 0.6804 - val_accuracy: 0.7576
Epoch 18/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5827 - accuracy:
```

```
racy: 0.7954 - val_loss: 0.6826 - val_accuracy: 0.7581
Epoch 19/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5823 - accuracy: 0.7958 - val_loss: 0.6818 - val_accuracy: 0.7580
Epoch 20/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.5820 - accuracy: 0.7953 - val_loss: 0.6835 - val_accuracy: 0.7579
```

model3: tfidf weighted fasttext embeddings

In [17]:

```
1 x_train_s1_tfidf_fasttext = np.load('train_s1_tfidf_fasttext.npy', mmap_mode='r')
2 x_train_s2_tfidf_fasttext = np.load('train_s2_tfidf_fasttext.npy', mmap_mode='r')
3 x_val_s1_tfidf_fasttext = np.load('val_s1_tfidf_fasttext.npy', mmap_mode='r')
4 x_val_s2_tfidf_fasttext = np.load('val_s2_tfidf_fasttext.npy', mmap_mode='r')
5 x_test_s1_tfidf_fasttext = np.load('test_s1_tfidf_fasttext.npy', mmap_mode='r')
6 x_test_s2_tfidf_fasttext = np.load('test_s2_tfidf_fasttext.npy', mmap_mode='r')
```

In [26]:

```
1 input2 = Input(shape=(300,)) #for sentence1 embeddings
2 norm2 = BatchNormalization()(input2)
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 norm3 = BatchNormalization()(input3)
5 merged = concatenate([norm2, norm3]) #merging all inputs
6 #norm1 = BatchNormalization()(merged)
7 dense1 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.01))
8 drop1 = Dropout(0.01)(dense1)
9 dense2 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.01))
10 drop2 = Dropout(0.01)(dense2)
11 dense3 = Dense(units = 100, activation = 'relu', kernel_regularizer = L1(0.01))
12 dense4 = Dense(units = 50, activation = 'relu')(dense3)
13 output = Dense(units = 3, activation = 'softmax')(dense4)
14 model3 = Model(inputs = [input2, input3], outputs = output)
```

In []:

```
1 model3.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics=['accuracy'])
2
3 tf.debugging.set_log_device_placement(True)
4
5 # Place tensors on the CPU
6 with tf.device('/CPU:0'):
7     model3.fit([x_train_s1_tfidf_fasttext, x_train_s2_tfidf_fasttext], y_train,
8                validation_data = ([x_val_s1_tfidf_fasttext, x_val_s2_tfidf_fasttext], y_val))
```

In [28]:

```
1 from keras import backend as K
2 K.set_value(model2_a.optimizer.learning_rate, 0.0001)
```

In [29]:

```
1 with tf.device('/CPU:0'):
2     model3.fit([x_train_s1_tfidf_fasttext, x_train_s2_tfidf_fasttext], y_tr
3                 validation_data = ([x_val_s1_tfidf_fasttext, x_val_s2_tfidf_fastt
Epoch 1/20
4292/4292 [=====] - 22s 5ms/step - loss: 0.6575 - accuracy: 0.7645 - val_loss: 0.7161 - val_accuracy: 0.7389
Epoch 2/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6574 - accuracy: 0.7648 - val_loss: 0.7147 - val_accuracy: 0.7398
Epoch 3/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6570 - accuracy: 0.7644 - val_loss: 0.7154 - val_accuracy: 0.7404
Epoch 4/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6566 - accuracy: 0.7650 - val_loss: 0.7211 - val_accuracy: 0.7345
Epoch 5/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6569 - accuracy: 0.7642 - val_loss: 0.7158 - val_accuracy: 0.7375
Epoch 6/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6562 - accuracy: 0.7650 - val_loss: 0.7148 - val_accuracy: 0.7400
Epoch 7/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6573 - accuracy: 0.7646 - val_loss: 0.7190 - val_accuracy: 0.7347
Epoch 8/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6566 - accuracy: 0.7646 - val_loss: 0.7182 - val_accuracy: 0.7392
Epoch 9/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6564 - accuracy: 0.7649 - val_loss: 0.7261 - val_accuracy: 0.7327
Epoch 10/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6561 - accuracy: 0.7654 - val_loss: 0.7198 - val_accuracy: 0.7415
Epoch 11/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6568 - accuracy: 0.7655 - val_loss: 0.7195 - val_accuracy: 0.7376
Epoch 12/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6566 - accuracy: 0.7652 - val_loss: 0.7172 - val_accuracy: 0.7385
Epoch 13/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6557 - accuracy: 0.7653 - val_loss: 0.7186 - val_accuracy: 0.7360
Epoch 14/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6568 - accuracy: 0.7652 - val_loss: 0.7143 - val_accuracy: 0.7430
Epoch 15/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6558 - accuracy: 0.7658 - val_loss: 0.7180 - val_accuracy: 0.7387
Epoch 16/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6566 - accuracy: 0.7650 - val_loss: 0.7182 - val_accuracy: 0.7397
Epoch 17/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6563 - accuracy: 0.7648 - val_loss: 0.7168 - val_accuracy: 0.7434
Epoch 18/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6564 - accu
```

```
racy: 0.7656 - val_loss: 0.7197 - val_accuracy: 0.7402
Epoch 19/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6566 - accuracy: 0.7650 - val_loss: 0.7179 - val_accuracy: 0.7381
Epoch 20/20
4292/4292 [=====] - 20s 5ms/step - loss: 0.6559 - accuracy: 0.7651 - val_loss: 0.7208 - val_accuracy: 0.7375
```

model4: tfidf weighted fasttext embeddings + extracted features

In [11]:

```
1 input1 = Input(shape=(29,)) #for extracted features
2 input2 = Input(shape=(300,)) #for sentence1 embeddings
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 merged = concatenate([input1, input2, input3]) #merging all inputs
5 norm1 = BatchNormalization()(merged)
6 dense1 = Dense(units = 100, activation = 'relu')(norm1)
7 dense2 = Dense(units = 80, activation = 'relu')(dense1)
8 dense3 = Dense(units = 50, activation = 'relu')(dense2)
9 dense4 = Dense(units = 40, activation = 'relu')(dense3)
10 dense5 = Dense(units = 30, activation = 'relu')(dense4)
11 dense6 = Dense(units = 20, activation = 'relu')(dense5)
12 dense7 = Dense(units = 10, activation = 'relu')(dense6)
13 output = Dense(units = 3, activation = 'softmax')(dense7)
14 model4 = Model(inputs = [input1, input2, input3], outputs = output)
```

In [18]:

```
1 model4.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
2 tf.debugging.set_log_device_placement(True)
# Place tensors on the CPU
4 with tf.device('/CPU:0'):
5     model4.fit([X_train_1, x_train_s1_tfidf_fasttext, x_train_s2_tfidf_fasttext],
6                 validation_data = ([X_val_1, x_val_s1_tfidf_fasttext, x_val_s2_tfidf_fasttext]),
6                 epochs=100)
7
8 Epoch 00/100
9 4292/4292 [=====] - 17s 4ms/step - loss: 0.4808 - accuracy: 0.8075 - val_loss: 0.6431 - val_accuracy: 0.7488
10 Epoch 87/100
11 4292/4292 [=====] - 18s 4ms/step - loss: 0.4800 - accuracy: 0.8078 - val_loss: 0.6495 - val_accuracy: 0.7484
12 Epoch 88/100
13 4292/4292 [=====] - 18s 4ms/step - loss: 0.4798 - accuracy: 0.8079 - val_loss: 0.6540 - val_accuracy: 0.7456
14 Epoch 89/100
15 4292/4292 [=====] - 17s 4ms/step - loss: 0.4797 - accuracy: 0.8081 - val_loss: 0.6521 - val_accuracy: 0.7466
16 Epoch 90/100
17 4292/4292 [=====] - 17s 4ms/step - loss: 0.4792 - accuracy: 0.8081 - val_loss: 0.6741 - val_accuracy: 0.7452
18 Epoch 91/100
19 4292/4292 [=====] - 17s 4ms/step - loss: 0.4786 - accuracy: 0.8088 - val_loss: 0.6479 - val_accuracy: 0.7469
20 Epoch 92/100
21 4292/4292 [=====] - 17s 4ms/step - loss: 0.4787 - accuracy: 0.8085 - val_loss: 0.6471 - val_accuracy: 0.7420
```

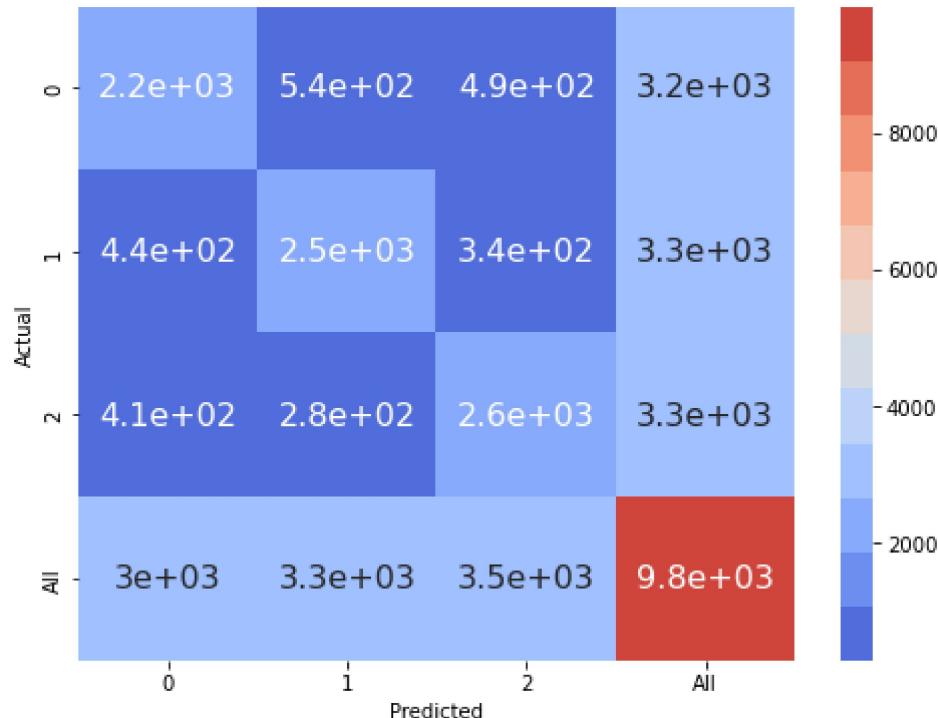
Error analysis on model4

train accuracy = 0.8, validation accuracy = 0.75

Confusion matrix: model4

```
In [19]: 1 val_pred = model4.predict([X_val_1, x_val_s1_tfidf_fasttext, x_val_s2_tfidf_])
In [22]: 1 label_pred = tf.argmax(val_pred, axis = 1)
In [31]: 1 data = {'y_Actual': y_val,
           2         'y_Predicted': label_pred}
           3 df = pd.DataFrame(data, columns=['y_Actual','y_Predicted'])
           4 confusion_matrix = pd.crosstab(df['y_Actual'], df['y_Predicted'], rownames=[

In [42]: 1 plt.figure(figsize = (8,6))
           2 colormap = sns.color_palette("coolwarm", 12)
           3 sns.heatmap(confusion_matrix, annot=True, cmap = colormap, annot_kws={"fontsize": 10})
           4 plt.show()
```



meaning of labels = 0:neutral, 1:contradiction, 2:entailment

- It seems that most errors are occurring for neutral pairs which are predicted as contradiction
- best performance is for entailment pairs
-

Erroneous points: model4

```
In [50]: 1 df2 = val_df[['sentence1', 'sentence2']]
```

```
In [51]: 1 df2['y_true'] = y_val
2 df2['y_pred'] = label_pred
```

C:\Users\anike\.conda\envs\tf_test\lib\site-packages\ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.

C:\Users\anike\.conda\envs\tf_test\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
In [53]: 1 # all incorrectly predicted pairs
2 erroneous_pairs = df2[df2['y_true']!=df2['y_pred']]
```

In [69]:

```
1 #true:entailment / pred:contradiction
2 ec = df2[(df2['y_true'] == 2) & (df2['y_pred'] == 1)]
3 print('some pairs of sentences belonging to entailment but predicted as cont')
4 print('.'*100)
5 for i,row in ec.sample(5).iterrows():
6     print('sentence1 : {}\\nsentence2 : {}\\nactual label : {}\\npredicted labe'
7         print('.'*100)
```

some pairs of sentences belonging to entailment but predicted as contradiction:

.....

sentence1 : A guy wearing a black shirt and gray pants pushing his shopping cart past a candy aisle of a Walmart.

sentence2 : A man is inside a store.

actual label : 2

predicted label : 1

.....

sentence1 : A girl is throwing a football on the beach.

sentence2 : A girl is playing with a football.

actual label : 2

predicted label : 1

.....

sentence1 : A middle-aged man in a bicycle race trying his hardest on the sand surrounded by watching crowd.

sentence2 : A man is riding a bicycle.

actual label : 2

predicted label : 1

.....

sentence1 : A motorcycle rider is making a sharp turn.

sentence2 : The motorcyclist is turning.

actual label : 2

predicted label : 1

.....

sentence1 : Men and women in swimsuits hangout on rocks above water.

sentence2 : The men and women are dressed in clothes for the beach.

actual label : 2

predicted label : 1

.....

Observations: It seems that the model is not being able to draw the meaning of the sentences.

Because two sentences here mean essentially the same ,still the model thinks that they are opposite in meaning. The model is incapable of learning the meaning, sequence information and symantics.

In [70]:

```
1 #true:contradiction / pred:contradiction
2 cc = df2[(df2['y_true'] == 1) & (df2['y_pred'] == 1)]
3 print('some pairs of sentences belonging to contradiction and predicted as c
4 print('.'*100)
5 for i,row in cc.sample(5).iterrows():
6     print('sentence1 : {}\\nsentence2 : {}\\nactual label : {}\\npredicted labe
7     print('.'*100)
```

some pairs of sentences belonging to contradiction and predicted as contradiction:

sentence1 : A boy wearing blue jeans is skateboarding.

sentence2 : a boy sleeps in a car

actual label : 1

predicted label : 1

sentence1 : A person in gray snowboarding down a hill.

sentence2 : The snowboarder is wearing white.

actual label : 1

predicted label : 1

sentence1 : A man in a green safety jacket and pants stands, facing away, near bags of trash while people stand on the curb across the street.

sentence2 : Nobody has a jacket

actual label : 1

predicted label : 1

sentence1 : A cyclists rides up a hill amid cheering bystanders.

sentence2 : The person is on a motorcycle.

actual label : 1

predicted label : 1

sentence1 : A man snowboards.

sentence2 : A man is drinking coffee at the cafe.

actual label : 1

predicted label : 1

In [71]:

```
1 #true:entailment / pred:entailment
2 ee = df2[(df2['y_true'] == 2) & (df2['y_pred'] == 2)]
3 print('some pairs of sentences belonging to entailment and predicted as entailment')
4 print('.'*100)
5 for i,row in ee.sample(5).iterrows():
6     print('sentence1 : {}\\nsentence2 : {}\\nactual label : {}\\npredicted label : {}\\n'.format(row['sentence1'], row['sentence2'], row['actual_label'], row['predicted_label']))
7     print('.'*100)
```

some pairs of sentences belonging to entailment and predicted as entailment:

.....

sentence1 : Two people in winter clothing sliding down a snow-covered hill.
sentence2 : People are sledding down a hill.

actual label : 2
predicted label : 2

.....

sentence1 : Speakers in a room full of people sitting in chairs.
sentence2 : There are people indoors.

actual label : 2
predicted label : 2

.....

sentence1 : A man in a black coat eats a doughnut while other men look out of frame.
sentence2 : A man in a black coat eats a doughnut.

actual label : 2
predicted label : 2

.....

sentence1 : A group of teenagers are standing in front of some tents, most of them are holding skateboards.
sentence2 : A group of teenagers are holding skateboards

actual label : 2
predicted label : 2

.....

sentence1 : A woman is writing something on a post-it note which is hanging on a bulletin board with a lot of other post-it notes.
sentence2 : The woman is writing a note.

actual label : 2
predicted label : 2

observation: It seems that when two sentences share some words or pos then the model is able to identify them correctly

In [68]:

```

1 #true:contradiction / pred:entailment
2 contradict_entailment = df2[(df2['y_true'] == 1) & (df2['y_pred'] == 2)]
3 print('some pairs of sentences belonging to contradiction but predicted as e
4 print('.'*100)
5 for i,row in contradict_entailment.sample(5).iterrows():
6     print('sentence1 : {}\\nsentence2 : {}\\nactual label : {}\\npredicted labe
7     print('.'*100)

```

some pairs of sentences belonging to contradiction but predicted as entailment:

.....

sentence1 : the shadow silhouette of a woman standing near the water looking at
a large attraction on the other side.

sentence2 : She is in the water.

actual label : 1

predicted label : 2

.....

sentence1 : A woman prepares ingredients for a bowl of soup.

sentence2 : A soup bowl prepares a woman.

actual label : 1

predicted label : 2

.....

sentence1 : A girl in a red polka dot bikini is jumping off a sand dune.

sentence2 : A girl in a bikini is buried in the sand.

actual label : 1

predicted label : 2

.....

sentence1 : A husky and a black cat nuzzling.

sentence2 : A cat attacks a dog.

actual label : 1

predicted label : 2

.....

sentence1 : A woman holds her child out of a red window, next to a Color TV sig
n.

sentence2 : The TV holds a woman who holds a child.

actual label : 1

predicted label : 2

.....

observation: Again we can see that the model is only considering presence of similar words to predict the label. It is not considering semantic meaning or polarity. Also it is not being able to recognize structure and voice of sentences.

In []:

1

