

In [28]:

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import os
6 import re
7 from nltk.corpus import stopwords
8 import pickle
9 from tqdm import tqdm
10 from wordcloud import WordCloud
11 from sklearn.feature_extraction.text import TfidfVectorizer
12 import tensorflow as tf
13 import random
14 import math
15 from textblob import TextBlob
16 from collections import Counter
17 from sklearn.metrics.pairwise import cosine_similarity
18 from sklearn.feature_extraction.text import TfidfVectorizer
19 import nltk, string
20 from collections import Counter
21 from nltk.stem import WordNetLemmatizer
22 import Levenshtein
23 import fasttext
24 import tensorflow as tf
25 from tensorflow.keras.layers import Embedding, LSTM, Dense, Bidirectional, c
26 from tensorflow.keras.models import Model
27 from tensorflow.keras.preprocessing.text import Tokenizer
28 from tensorflow.keras.preprocessing.sequence import pad_sequences
29 from tensorflow.keras.utils import plot_model
30 from tensorflow.keras.utils import to_categorical
31 from tensorflow.keras.regularizers import L1, L2
32 from tensorflow.keras.optimizers import Adam, SGD, RMSprop
33 from tensorflow.keras.utils import plot_model
34 from tensorflow.keras import backend as K
35 from tensorflow.keras.regularizers import l2, l1, l1_l2
36 %matplotlib inline

```

In [2]:

```

1 import tensorflow as tf
2 tf.config.experimental.allow_growth = True

```

In [3]:

```

1 import logging
2 logging.getLogger('tensorflow').setLevel(logging.ERROR) # suppress warnings

```

In [4]:

```

1 # Setting seed
2 global_seed = 42
3 np.random.seed(42)
4 random.seed(42)
5 tf.random.set_seed(42)

```

```
In [5]: 1 # Loading the dataframes in which features have already been extracted
2 train_df = pd.read_csv("train_df.csv")
3 val_df = pd.read_csv("val_df.csv")
4 test_df = pd.read_csv("test_df.csv")
```

```
In [6]: 1 # getting lists of sentence1 and sentence 2 for train and validation datasets
2 train_sentence1 = train_df['sentence1_preprocessed'].tolist()
3 train_sentence2 = train_df['sentence2_preprocessed'].tolist()
4 val_sentence1 = val_df['sentence1_preprocessed'].tolist()
5 val_sentence2 = val_df['sentence2_preprocessed'].tolist()
```

MLP on avg fasttext embeddings and extracted features

```
In [7]: 1 x_train_s1_avg_fasttext = np.load('x_train_s1_avg_fasttext.npy', mmap_mode='r')
2 x_train_s2_avg_fasttext = np.load('x_train_s2_avg_fasttext.npy', mmap_mode='r')
3 x_val_s1_avg_fasttext = np.load('x_val_s1_avg_fasttext.npy', mmap_mode='r')
4 x_val_s2_avg_fasttext = np.load('x_val_s2_avg_fasttext.npy', mmap_mode='r')
5 x_test_s1_avg_fasttext = np.load('x_test_s1_avg_fasttext.npy', mmap_mode='r')
6 x_test_s2_avg_fasttext = np.load('x_test_s2_avg_fasttext.npy', mmap_mode='r')
```

```
In [8]: 1 to_keep = ['log_bleu_score', 'sum_polarity', 'diff_subjectivity',
2             'cosine_similarity', 'tfidf_cosine_sim', 'diff_num_words', 's1_CC',
3             's1_IN', 's1_PRP', 's1_PRP$', 's1_RB', 's1_RBR', 's1_RBS', 's1_UH', 's
4             's2_PRP$', 's2_RB', 's2_RBR', 's2_RBS', 's2_UH', 'ratio_noun', 'ratio_
5             'diff_adjective', 'levenshtein_dist', 'tfidf_w2v_cosine',
6             'tfidf_cosine_sim_ngram']
```

```
In [9]: 1 X_train_1 = train_df[to_keep]
2 y_train = train_df['gold_label']
3 X_val_1 = val_df[to_keep]
4 y_val = val_df['gold_label']
```

```
In [10]: 1 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , va
2 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , valu
```

```
In [11]: 1 y_train_enc = to_categorical(y_train)
2 y_val_enc = to_categorical(y_val)
```

In [12]:

```
1 input1 = Input(shape=(29,)) #for extracted features
2 input2 = Input(shape=(300,)) #for sentence1 embeddings
3 input3 = Input(shape=(300,)) #for sentence2 embeddings
4 merged = concatenate([input1, input2, input3]) #merging all inputs
5 norm1 = BatchNormalization()(merged)
6 norm1 = tf.expand_dims(norm1, axis=-1)
7 conv1 = Conv1D(filters = 128, kernel_size = 3, strides=1, padding="valid")(n
8 conv2 = Conv1D(filters = 64, kernel_size = 3, strides=1, padding="valid")(co
9 conv3 = Conv1D(filters = 32, kernel_size = 3, strides=1, padding="valid")(co
10 flat = Flatten()(conv3)
11 dense1 = Dense(units = 1024, activation = 'relu')(flat)
12 drop1 = Dropout(0.3)(dense1)
13 dense2 = Dense(units = 512, activation = 'relu')(drop1)
14 drop2 = Dropout(0.3)(dense2)
15 dense3 = Dense(units = 256, activation = 'relu')(drop2)
16 drop3 = Dropout(0.3)(dense3)
17 norm2 = BatchNormalization()(drop3)
18 dense4 = Dense(units = 128, activation = 'relu')(norm2)
19 drop4 = Dropout(0.3)(dense4)
20 dense5 = Dense(units = 64, activation = 'relu')(dense4)
21 dense6 = Dense(units = 32, activation = 'relu')(dense5)
22 dense7 = Dense(units = 16, activation = 'relu')(dense6)
23
24 output = Dense(units = 3, activation = 'softmax')(dense7)
25 model1 = Model(inputs = [input1, input2, input3], outputs = output)
26
```

```
In [13]: 1 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
2 import tensorflow as tf
3 # Learning rate schedule
4 def step_decay(epoch):
5
6     import math
7     initial_lrate = 0.001
8     drop = 0.8
9     epochs_drop = 8.0
10    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
11    return lrate
12
13 lrate = LearningRateScheduler(step_decay)
14
15 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 10)
16 fname = 'model1.h5'
17 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save_
18 callbacks = [lrate,es]
19
20 model1.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric_
21 model1.fit([X_train_1, x_train_s1_avg_fasttext, x_train_s2_avg_fasttext], y_
22             validation_data = ([X_val_1, x_val_s1_avg_fasttext, x_val_s2_avg_
```

Epoch 1/50
8584/8584 [=====] - 138s 15ms/step - loss: 0.8259 -
accuracy: 0.6294 - val_loss: 0.7223 - val_accuracy: 0.6886
Epoch 2/50
8584/8584 [=====] - 133s 15ms/step - loss: 0.7248 -
accuracy: 0.6882 - val_loss: 0.6653 - val_accuracy: 0.7200
Epoch 3/50
8584/8584 [=====] - 134s 16ms/step - loss: 0.6826 -
accuracy: 0.7109 - val_loss: 0.6237 - val_accuracy: 0.7386
Epoch 4/50
8584/8584 [=====] - 137s 16ms/step - loss: 0.6552 -
accuracy: 0.7242 - val_loss: 0.6206 - val_accuracy: 0.7393
Epoch 5/50
8584/8584 [=====] - 136s 16ms/step - loss: 0.6347 -
accuracy: 0.7351 - val_loss: 0.6062 - val_accuracy: 0.7470
Epoch 6/50
8584/8584 [=====] - 137s 16ms/step - loss: 0.6181 -
accuracy: 0.7433 - val_loss: 0.5879 - val_accuracy: 0.7530
Epoch 7/50
8584/8584 [=====] - 138s 16ms/step - loss: 0.6015 -
accuracy: 0.7511 - val_loss: 0.5759 - val_accuracy: 0.7615

In []: 1

In []: 1

Sequence modeling

We will be using following general architecture for modelling:

```
In [15]: 1 from IPython.display import Image
          2 from IPython.core.display import HTML
          3 Image(url= "https://i.imgur.com/6q4df1V.png")
```

Out[15]:

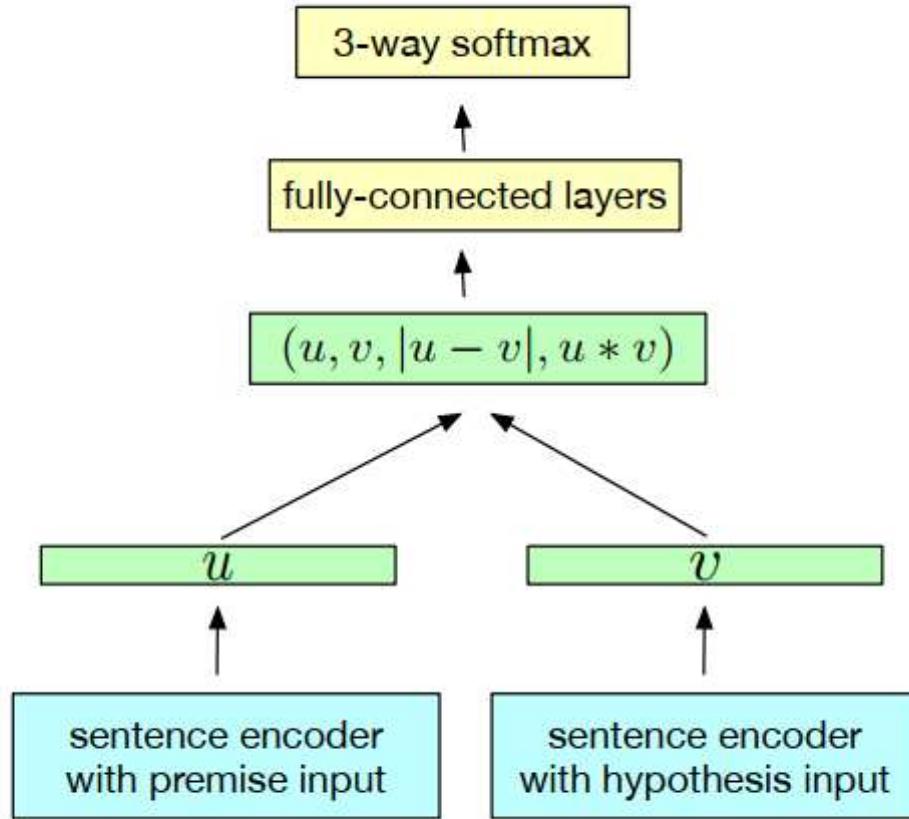


Figure 1: **Generic NLI training scheme.**

train sentence1 tokenize > token to ids > padding

In [12]:

```

1 num_words = None
2 oov_token = '<UNK>'
3 pad_type = 'post'
4 trunc_type = 'post'
5 # Tokenize our training data
6 tokenizer1 = Tokenizer(num_words=num_words, oov_token=oov_token)
7 tokenizer1.fit_on_texts(train_sentence1)
8
9 # Get our training data word index
10 word_index1 = tokenizer1.word_index
11
12 # Encode training data sentences into sequences
13 train_sequences_1 = tokenizer1.texts_to_sequences(train_sentence1)
14
15 # Get max training sequence Length
16 maxlen1 = max([len(x) for x in train_sequences_1])
17
18 # Pad the training sequences
19 train_padded_1 = pad_sequences(train_sequences_1, padding=pad_type, truncati

```

validation sentence1 > tokenize > token to ids > pad

In [13]:

```

1 # Encode training data sentences into sequences
2 val_sequences_1 = tokenizer1.texts_to_sequences(val_sentence1)
3
4 # Pad the training sequences
5 val_padded_1 = pad_sequences(val_sequences_1, padding=pad_type, truncating=t

```

train sentence2 tokenize > token to ids > padding

In [14]:

```

1 num_words = None
2 oov_token = '<UNK>'
3 pad_type = 'post'
4 trunc_type = 'post'
5 # Tokenize our training data
6 tokenizer2 = Tokenizer(num_words=num_words, oov_token=oov_token)
7 tokenizer2.fit_on_texts(train_sentence2)
8
9 # Get our training data word index
10 word_index2 = tokenizer2.word_index
11
12 # Encode training data sentences into sequences
13 train_sequences_2 = tokenizer2.texts_to_sequences(train_sentence2)
14
15 # Get max training sequence Length
16 maxlen2 = max([len(x) for x in train_sequences_2])
17
18 # Pad the training sequences
19 train_padded_2 = pad_sequences(train_sequences_2, padding=pad_type, truncati

```

validation sentence1 > tokenize > token to ids > pad

In [15]:

```

1 # Encode training data sentences into sequences
2 val_sequences_2 = tokenizer2.texts_to_sequences(val_sentence2)
3
4 # Pad the training sequences
5 val_padded_2 = pad_sequences(val_sequences_2, padding=pad_type, truncating=t

```

In [16]:

```

1 vocab_size_1 = len(word_index1) +1
2 vocab_size_2 = len(word_index2)+1

```

In [17]:

```

1 print(vocab_size_1)
2 print(vocab_size_2)

```

17997

30076

GloVe embeddings

Sentence1 tokens glove embeddings

In [18]:

```

1 #getting glove embeddings for tokens
2 embeddings_index_1 = dict()
3 f = open('glove.6B.300d.txt', encoding="utf8")
4 for line in f:
5     values = line.split()
6     word = values[0]
7     coefs = np.asarray(values[1:], dtype='float32')
8     embeddings_index_1[word] = coefs
9 f.close()
10
11 embedding_matrix_glove_1 = np.zeros((vocab_size_1, 300))
12 for word, i in word_index1.items():
13     embedding_vector = embeddings_index_1.get(word)
14     if embedding_vector is not None:
15         embedding_matrix_glove_1[i] = embedding_vector

```

Sentence2 tokens glove embeddings

In [19]:

```

1 embedding_matrix_glove_2 = np.zeros((vocab_size_2, 300))
2 for word, i in word_index2.items():
3     embedding_vector = embeddings_index_1.get(word)
4     if embedding_vector is not None:
5         embedding_matrix_glove_2[i] = embedding_vector

```

In [20]:

```

1 print(embedding_matrix_glove_1.shape)
2 print(embedding_matrix_glove_2.shape)

```

(17997, 300)
(30076, 300)

Fasttext embeddings

In [25]:

```

1 #Loading oretrained fasttext model
2 #https://fasttext.cc/docs/en/crawl-vectors.html
3 import fasttext.util
4 fasttext.util.download_model('en', if_exists='ignore') # English
5 ft = fasttext.load_model('cc.en.300.bin')

```

Sentence1 tokens fasttext embeddings

In [26]:

```

1 embedding_matrix_fasttext_1 = np.zeros((vocab_size_1, 300))
2 for word, i in word_index1.items():
3     embedding_vector = ft.get_word_vector(word)
4     if embedding_vector is not None:
5         embedding_matrix_fasttext_1[i] = embedding_vector

```

Sentence2 tokens fasttext embeddings

```
In [27]: 1 embedding_matrix_fasttext_2 = np.zeros((vocab_size_2, 300))
2 for word, i in word_index2.items():
3     embedding_vector = ft.get_word_vector(word)
4     if embedding_vector is not None:
5         embedding_matrix_fasttext_2[i] = embedding_vector
```

Encoding target variable

```
In [28]: 1 y_train = train_df['gold_label']
2 y_val = val_df['gold_label']
```

```
In [ ]: 1
```

```
In [ ]: 1 #converting target categories to numbers: neutral:0, contradiction:1, entailment:2
2 y_train.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , value=0, inplace=True)
3 y_val.replace(to_replace = ['neutral', 'contradiction', 'entailment'] , value=0, inplace=True)
```

```
In [31]: 1 #one hot encoding target
2 y_train_enc = to_categorical(y_train)
3 y_val_enc = to_categorical(y_val)
```

LSTM with glove embeddings

we will be taking max over each time step of LSTM

In [47]:

```

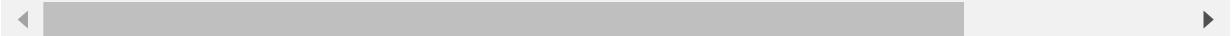
1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length =
6 weights = [embedding_matrix_glove_1], trainable = False)(in
7 lstm_s1 = LSTM(128, return_sequences = True)(emb1)
8 #lstm1 = LSTM(128, return_sequences = True)(lstm_s1)
9 rs1 = tf.keras.layers.Reshape(target_shape=(maxlen1,128,1))(lstm_s1)
10 avg1 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalMaxPool1D())(rs
11 rs2 = tf.keras.layers.Reshape(target_shape=(maxlen1,))(avg1)
12 s1_dense = Dense(units = 300, activation = 'relu' )(rs2)
13
14
15 #getting sentence2 representations
16 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
17 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length =
18 weights = [embedding_matrix_glove_2], trainable = False)(in
19 lstm_s2 = LSTM(128, return_sequences = True)(emb2)
20 #lstm2 = LSTM(128, return_sequences = True)(lstm_s2)
21 rs3 = tf.keras.layers.Reshape(target_shape=(maxlen2,128,1))(lstm_s2)
22 avg2 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalAveragePooling1
23 rs4 = tf.keras.layers.Reshape(target_shape=(maxlen2,))(avg2)
24 s2_dense = Dense(units = 300, activation = 'relu' )(rs4)
25
26
27 #merging sentence1 and sentence2 representations
28 merged_representations = concatenate([s1_dense, s2_dense])
29 prod = tf.math.multiply(s1_dense, s2_dense)
30 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
31 abs_diff = tf.math.abs(subtracted)
32
33 #concatenating all merged representations
34 concatenated = concatenate([merged_representations, prod, abs_diff])
35
36 #MLP
37 dense1 = Dense(units = 100, activation = 'relu')(concatenated)
38 drop1 = Dropout(0.3)(dense1)
39 dense2 = Dense(units = 50, activation = 'relu')(drop1)
40 drop2 = Dropout(0.3)(dense2)
41 dense3 = Dense(units = 25, activation = 'relu')(drop2)
42 drop3 = Dropout(0.3)(dense3)
43 output = Dense(units = 3, activation = 'softmax')(drop3)
44 model1 = Model(inputs = [input1, input2], outputs = output)
45
46

```

In [48]: 1 model1.summary()

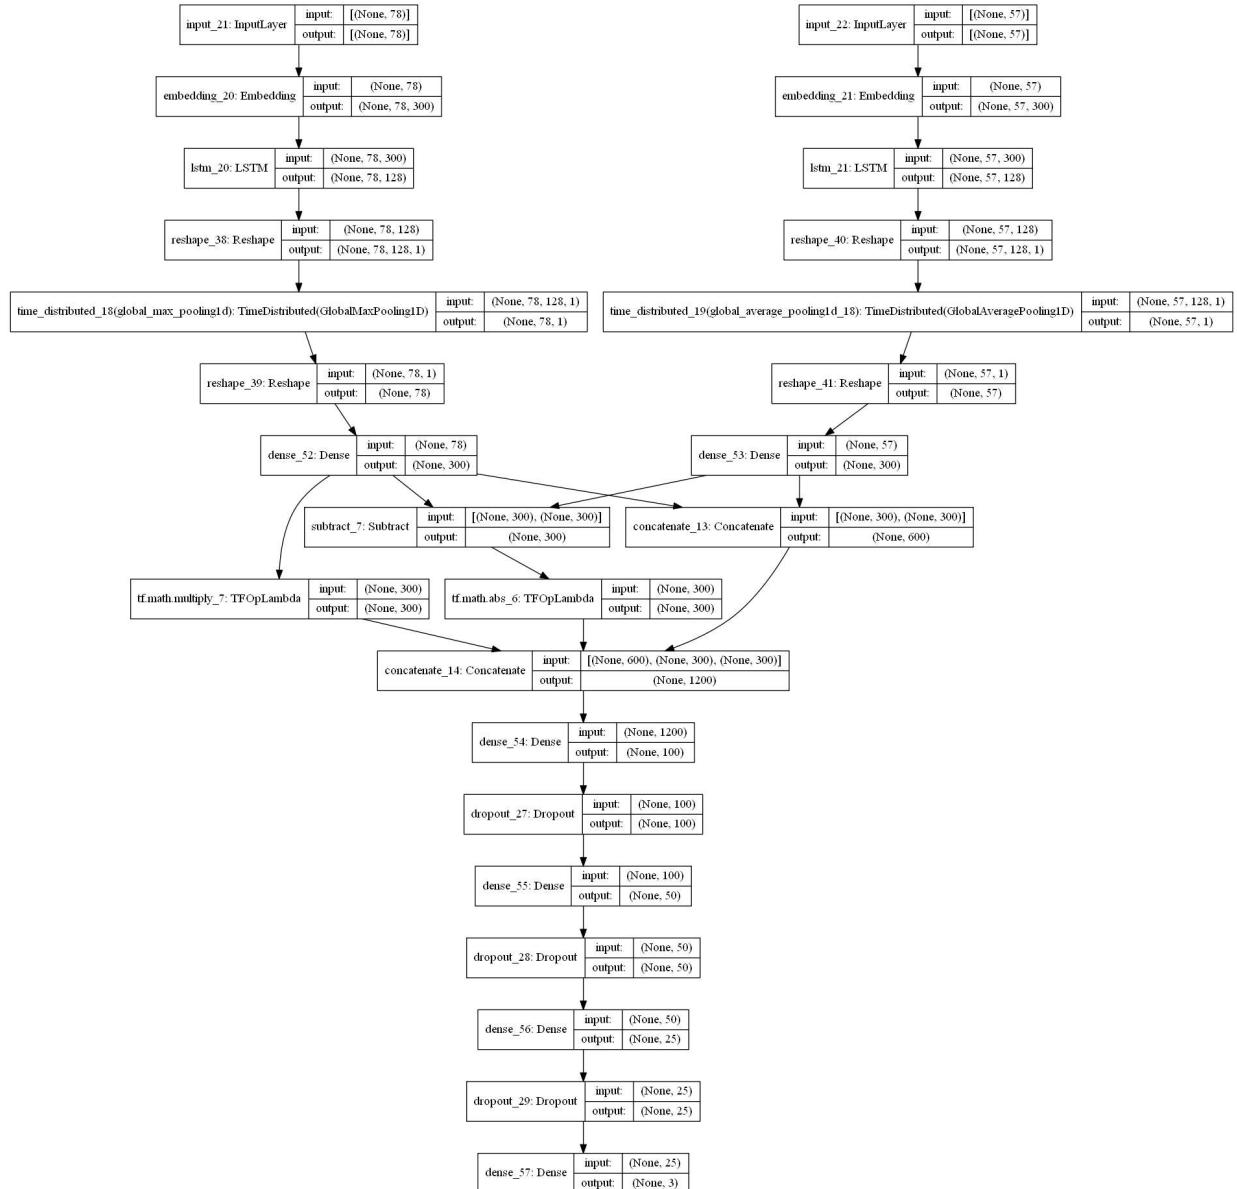
Model: "model_9"

| Layer (type) | Output Shape | Param # | Connected to |
|---------------------------------|--------------------|---------|---------------------------|
| input_21 (InputLayer) | [None, 78] | 0 | |
| embedding_20 (Embedding) | (None, 78, 300) | 5399100 | input_21[0][0] |
| embedding_21 (Embedding) | (None, 57, 300) | 9022800 | input_22[0][0] |
| lstm_20 (LSTM) | (None, 78, 128) | 219648 | embedding_20[0][0] |
| lstm_21 (LSTM) | (None, 57, 128) | 219648 | embedding_21[0][0] |
| reshape_38 (Reshape) | (None, 78, 128, 1) | 0 | lstm_20[0][0] |
| reshape_40 (Reshape) | (None, 57, 128, 1) | 0 | lstm_21[0][0] |
| time_distributed_18 (TimeDistr) | (None, 78, 1) | 0 | reshape_38[0][0] |
| time_distributed_19 (TimeDistr) | (None, 57, 1) | 0 | reshape_40[0][0] |
| reshape_39 (Reshape) | (None, 78) | 0 | time_distributed_18[0][0] |
| reshape_41 (Reshape) | (None, 57) | 0 | time_distributed_19[0][0] |
| dense_52 (Dense) | (None, 300) | 23700 | reshape_39[0][0] |
| dense_53 (Dense) | (None, 300) | 17400 | reshape_41[0][0] |

| | | | |
|--|--------------|--------|----------------------------------|
| <u>subtract_7</u> (Subtract) | (None, 300) | 0 | dense_52[0][0] dense_53[0][0] |
| <u>concatenate_13</u> (Concatenate) | (None, 600) | 0 | dense_52[0][0] dense_53[0][0] |
| <u>tf.math.multiply_7</u> (TFOpLambda) | (None, 300) | 0 | dense_52[0][0] dense_53[0][0] |
| <u>tf.math.abs_6</u> (TFOpLambda) | (None, 300) | 0 | subtract_7[0][0] |
| <u>concatenate_14</u> (Concatenate) | (None, 1200) | 0 | concatenate_13[0][0] |
| <u>tf.math.multiply_7</u> [0][0] | | | tf.math.multiply_7[0][0] |
| <u>tf.math.abs_6</u> [0][0] | | | tf.math.abs_6[0][0] |
| <u>dense_54</u> (Dense) | (None, 100) | 120100 | concatenate_14[0][0] |
| <u>dropout_27</u> (Dropout) | (None, 100) | 0 | dense_54[0][0] |
| <u>dense_55</u> (Dense) | (None, 50) | 5050 | dropout_27[0][0] |
| <u>dropout_28</u> (Dropout) | (None, 50) | 0 | dense_55[0][0] |
| <u>dense_56</u> (Dense) | (None, 25) | 1275 | dropout_28[0][0] |
| <u>dropout_29</u> (Dropout) | (None, 25) | 0 | dense_56[0][0] |
| <u>dense_57</u> (Dense) | (None, 3) | 78 | dropout_29[0][0] |
| <hr/> | | | |
| <hr/> | | | |
| Total params: 15,028,799 | | | |
| Trainable params: 606,899 | | | |
| Non-trainable params: 14,421,900 | | | |
| <hr/> | | | |
|  | | | |

In [49]: 1 plot_model(model1, to_file='model1.png', show_shapes=True)

Out[49]:



In [50]:

```
1 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
2 import tensorflow as tf
3 # Learning rate schedule
4 def step_decay(epoch):
5
6     import math
7     initial_lrate = 0.001
8     drop = 0.9
9     epochs_drop = 10.0
10    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
11    return lrate
12
13 lrate = LearningRateScheduler(step_decay)
14
15 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 5)
16 fname = 'best_model1.hdf5'
17 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save
18 callbacks = [lrate, es, checkpoint]
```

In [51]:

```
1 model1.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric
2 model1.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 50, batch_
3             validation_data = ([val_padded_1, val_padded_2], y_val_enc), call
```

```
Epoch 1/50
537/537 [=====] - 39s 65ms/step - loss: 0.9134 - accuracy: 0.5709 - val_loss: 0.8058 - val_accuracy: 0.6463
Epoch 2/50
537/537 [=====] - 35s 65ms/step - loss: 0.8029 - accuracy: 0.6473 - val_loss: 0.7561 - val_accuracy: 0.6712
Epoch 3/50
537/537 [=====] - 35s 66ms/step - loss: 0.7634 - accuracy: 0.6719 - val_loss: 0.7162 - val_accuracy: 0.6897
Epoch 4/50
537/537 [=====] - 36s 66ms/step - loss: 0.7351 - accuracy: 0.6884 - val_loss: 0.7070 - val_accuracy: 0.6974
Epoch 5/50
537/537 [=====] - 36s 67ms/step - loss: 0.7105 - accuracy: 0.7022 - val_loss: 0.6901 - val_accuracy: 0.7050
Epoch 6/50
537/537 [=====] - 36s 67ms/step - loss: 0.6918 - accuracy: 0.7124 - val_loss: 0.6741 - val_accuracy: 0.7141
Epoch 7/50
537/537 [=====] - 36s 67ms/step - loss: 0.6735 - accuracy: 0.7219 - val_loss: 0.6777 - val_accuracy: 0.7163
Epoch 8/50
537/537 [=====] - 36s 67ms/step - loss: 0.6583 - accuracy: 0.7300 - val_loss: 0.6552 - val_accuracy: 0.7258
Epoch 9/50
537/537 [=====] - 36s 67ms/step - loss: 0.6453 - accuracy: 0.7365 - val_loss: 0.6544 - val_accuracy: 0.7295
Epoch 10/50
537/537 [=====] - 36s 67ms/step - loss: 0.6304 - accuracy: 0.7435 - val_loss: 0.6642 - val_accuracy: 0.7297
Epoch 11/50
537/537 [=====] - 36s 67ms/step - loss: 0.6202 - accuracy: 0.7486 - val_loss: 0.6491 - val_accuracy: 0.7321
Epoch 12/50
537/537 [=====] - 36s 67ms/step - loss: 0.6105 - accuracy: 0.7535 - val_loss: 0.6457 - val_accuracy: 0.7327
Epoch 13/50
537/537 [=====] - 36s 67ms/step - loss: 0.6012 - accuracy: 0.7573 - val_loss: 0.6528 - val_accuracy: 0.7295
Epoch 14/50
537/537 [=====] - 36s 67ms/step - loss: 0.5917 - accuracy: 0.7619 - val_loss: 0.6603 - val_accuracy: 0.7330
Epoch 15/50
537/537 [=====] - 36s 67ms/step - loss: 0.5830 - accuracy: 0.7667 - val_loss: 0.6537 - val_accuracy: 0.7326
Epoch 16/50
537/537 [=====] - 36s 67ms/step - loss: 0.5754 - accuracy: 0.7699 - val_loss: 0.6627 - val_accuracy: 0.7337
Epoch 17/50
537/537 [=====] - 36s 67ms/step - loss: 0.5672 - accuracy: 0.7737 - val_loss: 0.6585 - val_accuracy: 0.7339
Epoch 18/50
537/537 [=====] - 36s 67ms/step - loss: 0.5610 - accuracy:
```

```
acy: 0.7770 - val_loss: 0.6708 - val_accuracy: 0.7333
Epoch 19/50
537/537 [=====] - 36s 67ms/step - loss: 0.5519 - accur
acy: 0.7809 - val_loss: 0.6789 - val_accuracy: 0.7317
Epoch 20/50
537/537 [=====] - 36s 67ms/step - loss: 0.5401 - accur
acy: 0.7870 - val_loss: 0.6895 - val_accuracy: 0.7291
Epoch 21/50
537/537 [=====] - 36s 67ms/step - loss: 0.5349 - accur
acy: 0.7891 - val_loss: 0.6818 - val_accuracy: 0.7301
Epoch 22/50
537/537 [=====] - 36s 67ms/step - loss: 0.5267 - accur
acy: 0.7932 - val_loss: 0.6844 - val_accuracy: 0.7318
```

Out[51]: <keras.callbacks.History at 0x1def70a6cf8>

LSTM model with fasttext embeddings

we will be taking max over each time step of LSTM

In [73]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length =
6   weights = [embedding_matrix_fasttext_1], trainable = False)
7 lstm_s1 = LSTM(128, return_sequences = True)(emb1)
8 print(lstm_s1.shape)
9 #lstm1 = LSTM(128, return_sequences = True)(lstm_s1)
10 rs1 = tf.keras.layers.Reshape(target_shape=(maxlen1,128,1))(lstm_s1)
11 avg1 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalMaxPool1D())(rs1)
12 rs2 = tf.keras.layers.Reshape(target_shape=(maxlen1,))(avg1)
13 s1_dense = Dense(units = 100, activation = 'relu' )(rs2)
14
15
16 #getting sentence2 representations
17 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
18 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length =
19   weights = [embedding_matrix_fasttext_2], trainable = False)
20 lstm_s2 = LSTM(128, return_sequences = True)(emb2)
21 #lstm2 = LSTM(128, return_sequences = True)(lstm_s2)
22 rs3 = tf.keras.layers.Reshape(target_shape=(maxlen2,128,1))(lstm_s2)
23 avg2 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalAveragePooling1D())(rs3)
24 rs4 = tf.keras.layers.Reshape(target_shape=(maxlen2,))(avg2)
25 s2_dense = Dense(units = 100, activation = 'relu' )(rs4)
26
27
28 #merging sentence1 and sentence2 representations
29 merged_representations = concatenate([s1_dense, s2_dense])
30 prod = tf.math.multiply(s1_dense, s2_dense)
31 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
32 abs_diff = tf.math.abs(subtracted)
33
34 #concatenating all merged representations
35 concatenated = concatenate([merged_representations, prod, abs_diff])
36
37 #MLP
38 dense1 = Dense(units = 100, activation = 'relu')(concatenated)
39 drop1 = Dropout(0.3)(dense1)
40 dense2 = Dense(units = 25, activation = 'relu')(drop1)
41 output = Dense(units = 3, activation = 'softmax')(dense2)
42 model2 = Model(inputs = [input1, input2], outputs = output)
43
44

```

(None, 78, 128)

In [68]: 1 model2.summary()

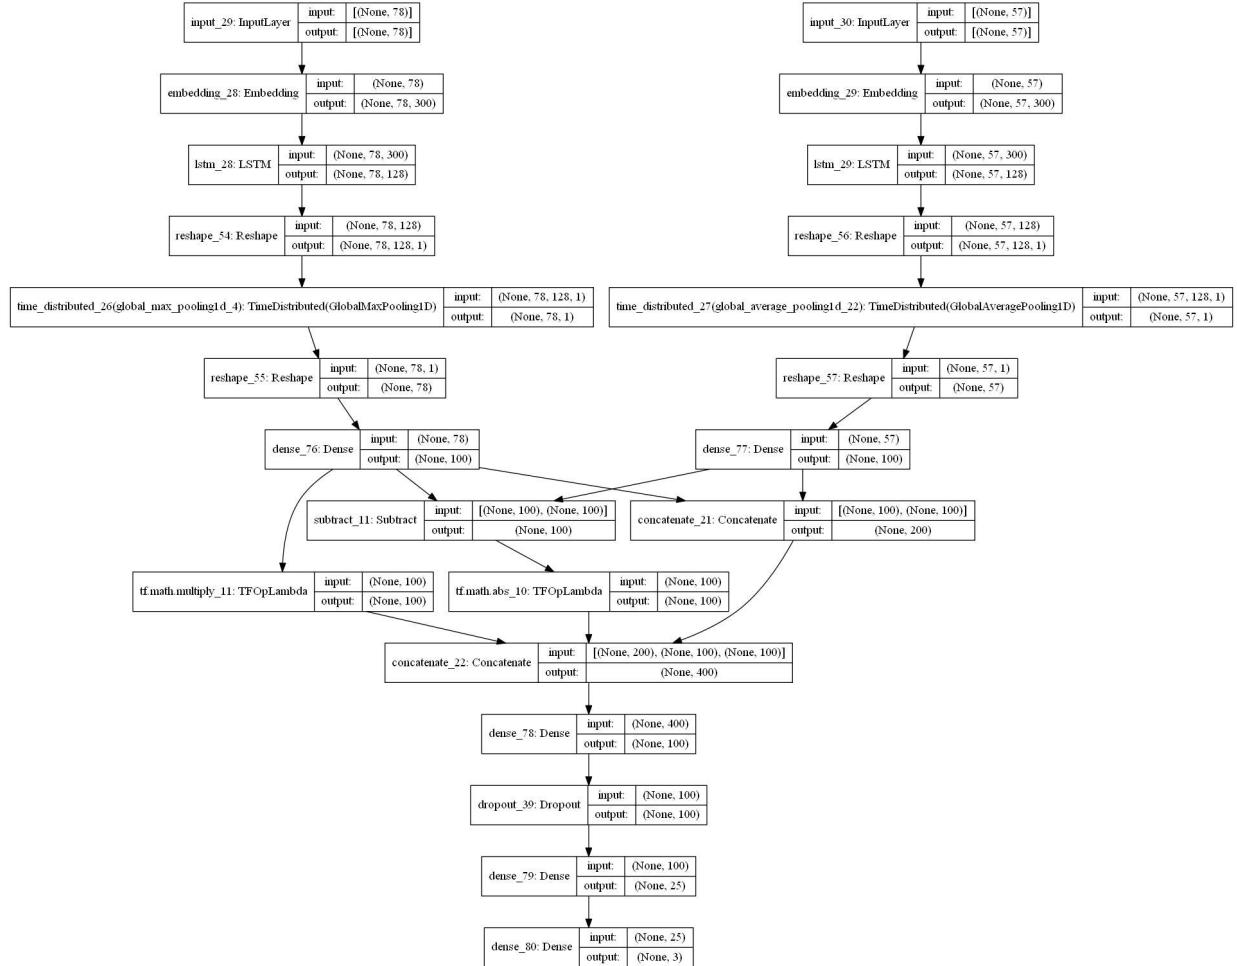
Model: "model_13"

| Layer (type) | Output Shape | Param # | Connected to |
|---------------------------------|--------------------|---------|---------------------------|
| input_29 (InputLayer) | [None, 78] | 0 | |
| embedding_28 (Embedding) | (None, 78, 300) | 5399100 | input_29[0][0] |
| embedding_29 (Embedding) | (None, 57, 300) | 9022800 | input_30[0][0] |
| lstm_28 (LSTM) | (None, 78, 128) | 219648 | embedding_28[0][0] |
| lstm_29 (LSTM) | (None, 57, 128) | 219648 | embedding_29[0][0] |
| reshape_54 (Reshape) | (None, 78, 128, 1) | 0 | lstm_28[0][0] |
| reshape_56 (Reshape) | (None, 57, 128, 1) | 0 | lstm_29[0][0] |
| time_distributed_26 (TimeDistr) | (None, 78, 1) | 0 | reshape_54[0][0] |
| time_distributed_27 (TimeDistr) | (None, 57, 1) | 0 | reshape_56[0][0] |
| reshape_55 (Reshape) | (None, 78) | 0 | time_distributed_26[0][0] |
| reshape_57 (Reshape) | (None, 57) | 0 | time_distributed_27[0][0] |
| dense_76 (Dense) | (None, 100) | 7900 | reshape_55[0][0] |
| dense_77 (Dense) | (None, 100) | 5800 | reshape_57[0][0] |

| | | | |
|--|-------------|-------|----------------------------------|
| <u>subtract_11</u> (Subtract) | (None, 100) | 0 | dense_76[0][0] dense_77[0][0] |
| <u>concatenate_21</u> (Concatenate) | (None, 200) | 0 | dense_76[0][0] dense_77[0][0] |
| <u>tf.math.multiply_11</u> (TFOpLambda (None, 100) | | 0 | dense_76[0][0] dense_77[0][0] |
| <u>tf.math.abs_10</u> (TFOpLambda [0] | (None, 100) | 0 | subtract_11[0] |
| <u>concatenate_22</u> (Concatenate) [0][0] | (None, 400) | 0 | concatenate_21 |
| <u>tf.math.multiply_11</u> [0][0] | | | tf.math.multiply_11[0][0] |
| <u>tf.math.abs_10</u> [0][0] | | | tf.math.abs_10 [0][0] |
| <u>dense_78</u> (Dense) [0][0] | (None, 100) | 40100 | concatenate_22 |
| <u>dropout_39</u> (Dropout) | (None, 100) | 0 | dense_78[0][0] |
| <u>dense_79</u> (Dense) [0] | (None, 25) | 2525 | dropout_39[0] |
| <u>dense_80</u> (Dense) | (None, 3) | 78 | dense_79[0][0] |
| <hr/> | | | |
| <hr/> | | | |
| Total params: 14,917,599 | | | |
| Trainable params: 495,699 | | | |
| Non-trainable params: 14,421,900 | | | |
| <hr/> | | | |
| <hr/> | | | |
|  | | | |
| <hr/> | | | |

In [69]: 1 plot_model(model2, to_file='model2.png', show_shapes=True)

Out[69]:



In [70]:

```

1 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
2 import tensorflow as tf
3 # Learning rate schedule
4 def step_decay(epoch):
5
6     import math
7     initial_lrate = 0.001
8     drop = 0.9
9     epochs_drop = 10.0
10    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
11    return lrate
12
13 lrate = LearningRateScheduler(step_decay)
14
15 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 5)
16 fname = 'best_model2.hdf5'
17 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save_best_only=True)
18 callbacks = [lrate, checkpoint]

```

In [71]:

```
1 model2.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric
2 model2.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 50, batch_
3 validation_data = ([val_padded_1, val_padded_2], y_val_enc), call
```

```
Epoch 1/50
537/537 [=====] - 39s 65ms/step - loss: 0.9331 - accuracy: 0.5515 - val_loss: 0.8487 - val_accuracy: 0.6084
Epoch 2/50
537/537 [=====] - 35s 65ms/step - loss: 0.8365 - accuracy: 0.6173 - val_loss: 0.8064 - val_accuracy: 0.6341
Epoch 3/50
537/537 [=====] - 35s 66ms/step - loss: 0.8049 - accuracy: 0.6353 - val_loss: 0.7946 - val_accuracy: 0.6428
Epoch 4/50
537/537 [=====] - 35s 66ms/step - loss: 0.7847 - accuracy: 0.6486 - val_loss: 0.7751 - val_accuracy: 0.6569
Epoch 5/50
537/537 [=====] - 35s 66ms/step - loss: 0.7614 - accuracy: 0.6643 - val_loss: 0.7468 - val_accuracy: 0.6719
Epoch 6/50
537/537 [=====] - 36s 66ms/step - loss: 0.7399 - accuracy: 0.6770 - val_loss: 0.7252 - val_accuracy: 0.6874
Epoch 7/50
537/537 [=====] - 36s 66ms/step - loss: 0.7200 - accuracy: 0.6900 - val_loss: 0.7052 - val_accuracy: 0.6974
```

Bidirectional LSTM with fasttext embedding

In [27]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length =
6                         weights = [embedding_matrix_fasttext_1], trainable = False)
7 lstm_s1 = Bidirectional(LSTM(200, return_sequences = True) , merge_mode = 'c'
8 lstm1 = Bidirectional(LSTM(100, return_sequences = True))(lstm_s1)
9 rs1 = tf.keras.layers.Reshape(target_shape=(maxlen1,200,1))(lstm1)
10 avg1 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalMaxPool1D())(rs1)
11 rs2 = tf.keras.layers.Reshape(target_shape=(maxlen1,))(avg1)
12 s1_dense = Dense(units = 100, activation = 'relu' )(rs2)
13
14
15 #getting sentence2 representations
16 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
17 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length =
18                         weights = [embedding_matrix_fasttext_2], trainable = False)
19 lstm_s2 = Bidirectional(LSTM(200, return_sequences = True))(emb2)
20 lstm2 = Bidirectional(LSTM(100, return_sequences = True))(lstm_s2)
21 rs3 = tf.keras.layers.Reshape(target_shape=(maxlen2,200,1))(lstm2)
22 avg2 = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalAveragePooling1D())
23 rs4 = tf.keras.layers.Reshape(target_shape=(maxlen2,))(avg2)
24 s2_dense = Dense(units = 100, activation = 'relu' )(rs4)
25
26
27 #merging sentence1 and sentence2 representations
28 merged_representations = concatenate([s1_dense, s2_dense])
29 prod = tf.math.multiply(s1_dense, s2_dense)
30 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
31 abs_diff = tf.math.abs(subtracted)
32
33 #concatenating all merged representations
34 concatenated = concatenate([merged_representations, prod, abs_diff])
35
36 #MLP
37 dense1 = Dense(units = 100, activation = 'relu')(concatenated)
38 drop1 = Dropout(0.3)(dense1)
39 dense2 = Dense(units = 50, activation = 'relu')(drop1)
40 drop2 = Dropout(0.2)(dense2)
41 dense3 = Dense(units = 10, activation = 'relu')(drop2)
42 output = Dense(units = 3, activation = 'softmax')(dense3)
43 model3 = Model(inputs = [input1, input2], outputs = output)
44
45

```

In [28]: 1 model3.summary()

Model: "model_3"

| Layer (type) | Output Shape | Param # | Connected to |
|--|--------------------|---------|-----------------------|
| input_7 (InputLayer) | [None, 78] | 0 | |
| input_8 (InputLayer) | [None, 57] | 0 | |
| embedding_6 (Embedding) | (None, 78, 300) | 5399100 | input_7[0][0] |
| embedding_7 (Embedding) | (None, 57, 300) | 9022800 | input_8[0][0] |
| bidirectional_12 (Bidirectional (None, 78, 400) [0]) | | 801600 | embedding_6[0] |
| bidirectional_14 (Bidirectional (None, 57, 400) [0]) | | 801600 | embedding_7[0] |
| bidirectional_13 (Bidirectional (None, 78, 200) [0]) | | 400800 | bidirectional_12[0] |
| bidirectional_15 (Bidirectional (None, 57, 200) [0]) | | 400800 | bidirectional_14[0] |
| reshape_12 (Reshape) | (None, 78, 200, 1) | 0 | bidirectional_13[0] |
| reshape_14 (Reshape) | (None, 57, 200, 1) | 0 | bidirectional_15[0] |
| time_distributed_6 (TimeDistrib (None, 78, 1) [0]) | | 0 | reshape_12[0] |
| time_distributed_7 (TimeDistrib (None, 57, 1) [0]) | | 0 | reshape_14[0] |
| reshape_13 (Reshape) | (None, 78) | 0 | time_distributed_6[0] |
| reshape_15 (Reshape) | (None, 57) | 0 | time_distributed_7[0] |

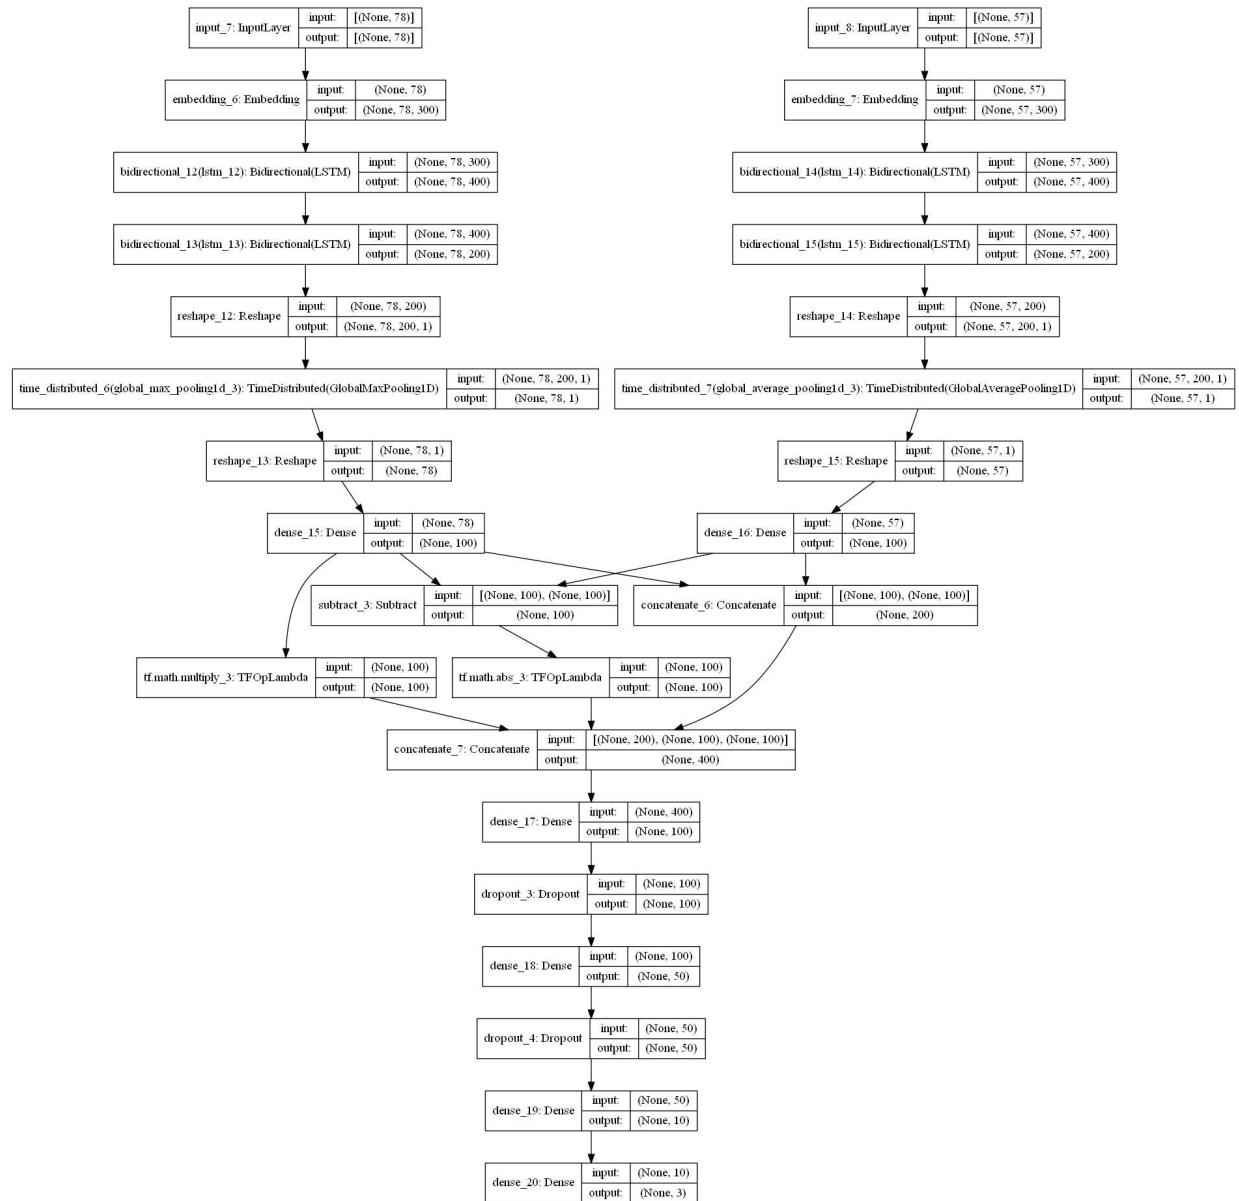
ed_7[0][0]

| | | | |
|---|-------------|-------|---|
| dense_15 (Dense) [0] | (None, 100) | 7900 | reshape_13[0] |
| dense_16 (Dense) [0] | (None, 100) | 5800 | reshape_15[0] |
| subtract_3 (Subtract) | (None, 100) | 0 | dense_15[0][0] dense_16[0][0] |
| concatenate_6 (Concatenate) | (None, 200) | 0 | dense_15[0][0] dense_16[0][0] |
| tf.math.multiply_3 (TFOpLambda) (None, 100) | 0 | 0 | dense_15[0][0] dense_16[0][0] |
| tf.math.abs_3 (TFOpLambda) [0] | (None, 100) | 0 | subtract_3[0] |
| concatenate_7 (Concatenate) [0][0] | (None, 400) | 0 | concatenate_6 tf.math.multiply_3[0][0] |
| tf.math.abs_3 [0][0] | | | |
| dense_17 (Dense) [0][0] | (None, 100) | 40100 | concatenate_7 |
| dropout_3 (Dropout) | (None, 100) | 0 | dense_17[0][0] |
| dense_18 (Dense) [0] | (None, 50) | 5050 | dropout_3[0] |
| dropout_4 (Dropout) | (None, 50) | 0 | dense_18[0][0] |
| dense_19 (Dense) [0] | (None, 10) | 510 | dropout_4[0] |
| dense_20 (Dense) | (None, 3) | 33 | dense_19[0][0] |
| <hr/> | | | |
| ===== | | | |
| ===== | | | |
| Total params: 16,886,093 | | | |

Trainable params: 2,464,193
 Non-trainable params: 14,421,900

In [29]: 1 plot_model(model3, to_file='model3.png', show_shapes=True)

Out[29]:



In [30]:

```
1 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
2 import tensorflow as tf
3 # Learning rate schedule
4 def step_decay(epoch):
5
6     import math
7     initial_lrate = 0.001
8     drop = 0.9
9     epochs_drop = 10.0
10    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
11    return lrate
12
13 lrate = LearningRateScheduler(step_decay)
14
15 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 6)
16 fname = 'best_model3.hdf5'
17 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save
18 callbacks = [es, lrate, checkpoint]
```

In [31]:

```

1 model3.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metric
2 model3.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 50, batch_
3 validation_data = ([val_padded_1, val_padded_2], y_val_enc), call

```

Epoch 1/50
4292/4292 [=====] - 327s 73ms/step - loss: 0.8976 - accuracy: 0.5820 - val_loss: 0.8098 - val_accuracy: 0.6443
Epoch 2/50
4292/4292 [=====] - 321s 75ms/step - loss: 0.8047 - accuracy: 0.6427 - val_loss: 0.7841 - val_accuracy: 0.6557
Epoch 3/50
4292/4292 [=====] - 323s 75ms/step - loss: 0.7770 - accuracy: 0.6584 - val_loss: 0.7618 - val_accuracy: 0.6654
Epoch 4/50
4292/4292 [=====] - 320s 74ms/step - loss: 0.7573 - accuracy: 0.6696 - val_loss: 0.7518 - val_accuracy: 0.6745
Epoch 5/50
4292/4292 [=====] - 323s 75ms/step - loss: 0.7414 - accuracy: 0.6775 - val_loss: 0.7477 - val_accuracy: 0.6746
Epoch 6/50
4292/4292 [=====] - 324s 75ms/step - loss: 0.7270 - accuracy: 0.6858 - val_loss: 0.7446 - val_accuracy: 0.6781
Epoch 7/50
4292/4292 [=====] - 324s 75ms/step - loss: 0.7127 - accuracy: 0.6929 - val_loss: 0.7415 - val_accuracy: 0.6776
Epoch 8/50
4292/4292 [=====] - 324s 75ms/step - loss: 0.6988 - accuracy: 0.6999 - val_loss: 0.7452 - val_accuracy: 0.6812
Epoch 9/50
4292/4292 [=====] - 324s 76ms/step - loss: 0.6834 - accuracy: 0.7079 - val_loss: 0.7467 - val_accuracy: 0.6822
Epoch 10/50
4292/4292 [=====] - 324s 76ms/step - loss: 0.6649 - accuracy: 0.7170 - val_loss: 0.7536 - val_accuracy: 0.6783
Epoch 11/50
4292/4292 [=====] - 340s 79ms/step - loss: 0.6477 - accuracy: 0.7254 - val_loss: 0.7567 - val_accuracy: 0.6794
Epoch 12/50
4292/4292 [=====] - 326s 76ms/step - loss: 0.6300 - accuracy: 0.7341 - val_loss: 0.7668 - val_accuracy: 0.6749
Epoch 13/50
4292/4292 [=====] - 326s 76ms/step - loss: 0.6121 - accuracy: 0.7429 - val_loss: 0.7776 - val_accuracy: 0.6730
Epoch 14/50
4292/4292 [=====] - 326s 76ms/step - loss: 0.5939 - accuracy: 0.7515 - val_loss: 0.8028 - val_accuracy: 0.6672
Epoch 15/50
4292/4292 [=====] - 325s 76ms/step - loss: 0.5749 - accuracy: 0.7604 - val_loss: 0.8195 - val_accuracy: 0.6703

Out[31]: <keras.callbacks.History at 0x214a0428978>

Bi-LSTM with attention with glove embeddings

```
In [32]: 1 import logging
          2 logging.getLogger('tensorflow').setLevel(logging.ERROR) # suppress warnings
```

Following self attention has been applied:

$$\begin{aligned}\overline{h}_i &= v_i \cdot \tanh(W \cdot h_i + b) + v_b \\ \alpha_i &= \text{softmax}(\overline{h}_i) \\ u &= \sum \alpha_i h_i\end{aligned}$$

where, h_i represent **Istm hidden states**, W and v represent affine transformations to calculate weights for each hidden states

```
In [24]: 1 class Attention_Custom(tf.keras.layers.Layer):
          2
          3     def __init__(self, att_units, inp_shape):
          4         super().__init__()
          5         self.att_units = att_units
          6         self.inp_shape = inp_shape
          7         self.num_dim = self.inp_shape[-1] ##dimensions per time step
          8         self.ts = self.inp_shape[-2] ##number of timesteps
          9         self.w = self.add_weight(shape = (self.num_dim, self.att_units), init
          10         self.b = self.add_weight(shape = (self.ts, self.att_units), initializ
          11         self.v = self.add_weight(shape = (self.att_units, self.num_dim), ini
          12         self.v_b = self.add_weight(shape = (self.ts, self.num_dim))
          13
          14     def call(self, x):
          15         score = K.dot((K.tanh(K.dot(x, self.w)+self.b)),self.v)+self.v_b
          16         weights = K.softmax(score, axis=1)
          17         weighted_states = x*weights
          18         output = K.sum(weighted_states, axis=1)
          19
          20     return weights, output
          21
```

In [40]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length =
6 weights = [embedding_matrix_glove_1], trainable = False)(in
7 #lstm_s1 = Bidirectional(LSTM(256, return_sequences = True) , merge_mode =
8 lstm1 = Bidirectional(LSTM(128, return_sequences = True, dropout=0.3, kernel_
9 merge_mode = 'concat')(emb1)
10 att_1 = Attention_Custom(64, lstm1.shape)
11 att_weights_1, context_1 = att_1.call(lstm1)
12 s1_dense = Dense(units = 300, activation = 'relu' )(context_1)
13
14
15 #getting sentence2 representations
16 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
17 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length =
18 weights = [embedding_matrix_glove_2], trainable = False)(in
19 #lstm_s2 = Bidirectional(LSTM(20, return_sequences = True))(emb2)
20 lstm2 = Bidirectional(LSTM(128, return_sequences = True, dropout=0.3, kernel_
21 merge_mode = 'concat')(emb2)
22 att_2 = Attention_Custom(64, lstm2.shape)
23 att_weights_2, context_2 = att_2.call(lstm2)
24 s2_dense = Dense(units = 300, activation = 'relu' )(context_2)
25
26 #merging sentence1 and sentence2 representations
27 merged_representations = concatenate([s1_dense, s2_dense])
28 #point-wise product of sentence representations
29 prod = tf.math.multiply(s1_dense, s2_dense)
30 #point-wise absolute difference of sentence representations
31 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
32 abs_diff = tf.math.abs(subtracted)
33
34 #concatenating
35 concatenated = concatenate([merged_representations, prod, abs_diff])
36 # concatenated = tf.expand_dims(concatenated, axis=-1)
37 # conv1 = Conv1D(filters = 64, kernel_size = 3, strides=1, padding='valid')
38 # conv2 = Conv1D(filters = 32, kernel_size = 3, strides=1, padding='valid')
39 # conv3 = Conv1D(filters = 16, kernel_size = 3, strides=1, padding='valid')
40
41 # flat = Flatten()(conv3)
42 #MLP and three-way classifier
43 dense1 = Dense(units = 512, activation = 'relu')(concatenated)
44 #drop1 = Dropout(0.3)(dense1)
45 # dense2 = Dense(units = 50, activation = 'relu')(drop1)
46 output = Dense(units = 3, activation = 'softmax')(dense1)
47 model17 = Model(inputs = [input1, input2], outputs = output)
48
49
50

```

```
In [41]: 1 model7.summary()
```

Model: "model_6"

| Layer (type) | Output Shape | Param # | Connected to |
|---------------------------------|-----------------|---------|--------------|
| input_13 (InputLayer) | [None, 78] | 0 | |
| input_14 (InputLayer) | [None, 57] | 0 | |
| embedding_12 (Embedding) [0] | (None, 78, 300) | 5399100 | input_13[0] |
| embedding_13 (Embedding) [0] | (None, 57, 300) | 9022800 | input_14[0] |

In [42]:

```

1 import tensorflow_addons as tfa
2 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
3 import tensorflow as tf
4 # Learning rate schedule
5 def step_decay(epoch):
6
7     import math
8     initial_lrate = 0.1
9     drop = 0.9
10    epochs_drop = 10.0
11    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
12    return lrate
13
14 #lrate = LearningRateScheduler(step_decay)
15 lrate = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1, patience=10, m
16 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 10)
17 fname = 'best_model7.hdf5'
18 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save_
19 callbacks = [lrate,es]
20 #optimizer=tfa.optimizers.SGDW(weight_decay=0.99,momentum=0.9, Learning_rate=
21 #optimizer=tfa.optimizers.AdamW(weight_decay=0.99, Learning_rate=0.001)
22 optimizer=tf.optimizers.Adam(learning_rate=0.001)
23
24
25 #sgd = SGD(Learning_rate = 0.1, )
26 model7.compile(optimizer = optimizer, loss = 'categorical_crossentropy', met
27 model7.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 100, batc
28             validation_data = ([val_padded_1, val_padded_2], y_val_enc), call

```

Epoch 1/100
8584/8584 [=====] - 337s 39ms/step - loss: 0.7877 - accuracy: 0.6600 - val_loss: 0.6641 - val_accuracy: 0.7384
Epoch 2/100
8584/8584 [=====] - 339s 39ms/step - loss: 0.6718 - accuracy: 0.7328 - val_loss: 0.6226 - val_accuracy: 0.7617
Epoch 3/100
8584/8584 [=====] - 334s 39ms/step - loss: 0.6368 - accuracy: 0.7543 - val_loss: 0.5997 - val_accuracy: 0.7712
Epoch 4/100
8584/8584 [=====] - 341s 40ms/step - loss: 0.6154 - accuracy: 0.7662 - val_loss: 0.5864 - val_accuracy: 0.7770
Epoch 5/100
8584/8584 [=====] - 341s 40ms/step - loss: 0.5994 - accuracy: 0.7760 - val_loss: 0.5791 - val_accuracy: 0.7908
Epoch 6/100
8584/8584 [=====] - 356s 41ms/step - loss: 0.5876 - accuracy: 0.7830 - val_loss: 0.5771 - val_accuracy: 0.7918
Epoch 7/100
8584/8584 [=====] - 341s 40ms/step - loss: 0.5776 - accuracy: 0.7884 - val_loss: 0.5724 - val_accuracy: 0.7929
Epoch 8/100
8584/8584 [=====] - 334s 39ms/step - loss: 0.5691 - accuracy: 0.7937 - val_loss: 0.5616 - val_accuracy: 0.7991
Epoch 9/100
8584/8584 [=====] - 333s 39ms/step - loss: 0.5621 - accuracy: 0.7977 - val_loss: 0.5693 - val_accuracy: 0.7982
Epoch 10/100

```
8584/8584 [=====] - 334s 39ms/step - loss: 0.5554 - accuracy: 0.8014 - val_loss: 0.5792 - val_accuracy: 0.7997
Epoch 11/100
8584/8584 [=====] - 332s 39ms/step - loss: 0.5498 - accuracy: 0.8049 - val_loss: 0.5648 - val_accuracy: 0.8038
Epoch 12/100
8584/8584 [=====] - 340s 40ms/step - loss: 0.5440 - accuracy: 0.8079 - val_loss: 0.5774 - val_accuracy: 0.7994
Epoch 13/100
8584/8584 [=====] - 339s 39ms/step - loss: 0.5396 - accuracy: 0.8104 - val_loss: 0.5761 - val_accuracy: 0.8025
Epoch 14/100
8584/8584 [=====] - 334s 39ms/step - loss: 0.5349 - accuracy: 0.8127 - val_loss: 0.5764 - val_accuracy: 0.8041
Epoch 15/100
8584/8584 [=====] - 331s 39ms/step - loss: 0.5312 - accuracy: 0.8146 - val_loss: 0.5680 - val_accuracy: 0.8068
Epoch 16/100
8584/8584 [=====] - 332s 39ms/step - loss: 0.5263 - accuracy: 0.8172 - val_loss: 0.5717 - val_accuracy: 0.8104
Epoch 17/100
8584/8584 [=====] - 335s 39ms/step - loss: 0.5231 - accuracy: 0.8194 - val_loss: 0.5851 - val_accuracy: 0.8022
Epoch 18/100
8584/8584 [=====] - 336s 39ms/step - loss: 0.5194 - accuracy: 0.8213 - val_loss: 0.5743 - val_accuracy: 0.8074
Epoch 19/100
8584/8584 [=====] - 328s 38ms/step - loss: 0.5156 - accuracy: 0.8231 - val_loss: 0.5862 - val_accuracy: 0.8032
Epoch 20/100
8584/8584 [=====] - 332s 39ms/step - loss: 0.5117 - accuracy: 0.8248 - val_loss: 0.5768 - val_accuracy: 0.8093
Epoch 21/100
8584/8584 [=====] - 323s 38ms/step - loss: 0.5088 - accuracy: 0.8266 - val_loss: 0.5843 - val_accuracy: 0.8068
Epoch 22/100
8584/8584 [=====] - 322s 37ms/step - loss: 0.5067 - accuracy: 0.8275 - val_loss: 0.5950 - val_accuracy: 0.8028
Epoch 23/100
8584/8584 [=====] - 321s 37ms/step - loss: 0.5036 - accuracy: 0.8299 - val_loss: 0.5988 - val_accuracy: 0.8024
Epoch 24/100
8584/8584 [=====] - 321s 37ms/step - loss: 0.5011 - accuracy: 0.8307 - val_loss: 0.5991 - val_accuracy: 0.8079
Epoch 25/100
8584/8584 [=====] - 322s 37ms/step - loss: 0.4988 - accuracy: 0.8317 - val_loss: 0.5854 - val_accuracy: 0.8067
Epoch 26/100
8584/8584 [=====] - 321s 37ms/step - loss: 0.4957 - accuracy: 0.8338 - val_loss: 0.6001 - val_accuracy: 0.8017
```

Out[42]: <keras.callbacks.History at 0x1f54c56d438>

Bi-LSTM with Attention on fasttext embeddings

In [34]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length=
6                         weights = [embedding_matrix_fasttext_1], trainable = False)
7 #lstm_s1 = Bidirectional(LSTM(256, return_sequences = True) , merge_mode =
8 lstm1 = Bidirectional(LSTM(128, return_sequences = True, dropout=0.3, kernel_
9                         merge_mode = 'concat'))(emb1)
10 att_1 = Attention_Custom(64, lstm1.shape)
11 att_weights_1, context_1 = att_1.call(lstm1)
12 s1_dense = Dense(units = 300, activation = 'relu' )(context_1)
13
14
15 #getting sentence2 representations
16 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
17 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length=
18                         weights = [embedding_matrix_fasttext_2], trainable = False)
19 #lstm_s2 = Bidirectional(LSTM(20, return_sequences = True))(emb2)
20 lstm2 = Bidirectional(LSTM(128, return_sequences = True, dropout=0.3, kernel_
21                         merge_mode = 'concat'))(emb2)
22 att_2 = Attention_Custom(64, lstm2.shape)
23 att_weights_2, context_2 = att_2.call(lstm2)
24 s2_dense = Dense(units = 300, activation = 'relu' )(context_2)
25
26 #merging sentence1 and sentence2 representations
27 merged_representations = concatenate([s1_dense, s2_dense])
28 #point-wise product of sentence representations
29 prod = tf.math.multiply(s1_dense, s2_dense)
30 #point-wise absolute difference of sentence representations
31 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
32 abs_diff = tf.math.abs(subtracted)
33
34 #concatenating
35 concatenated = concatenate([merged_representations, prod, abs_diff])
36 # concatenated = tf.expand_dims(concatenated, axis=-1)
37 # conv1 = Conv1D(filters = 64, kernel_size = 3, strides=1, padding='valid')(
38 # conv2 = Conv1D(filters = 32, kernel_size = 3, strides=1, padding='valid')(
39 # conv3 = Conv1D(filters = 16, kernel_size = 3, strides=1, padding='valid')(
40
41 # flat = Flatten()(conv3)
42 #MLP and three-way classifier
43 dense1 = Dense(units = 512, activation = 'relu')(concatenated)
44 drop1 = Dropout(0.3)(dense1)
45 # dense2 = Dense(units = 50, activation = 'relu')(drop1)
46 output = Dense(units = 3, activation = 'softmax')(dense1)
47 model8 = Model(inputs = [input1, input2], outputs = output)
48
49
50

```

```
In [35]: 1 model8.summary()
```

Model: "model_1"

| Layer (type) | Output Shape | Param # | Connected to |
|--------------------------------|-----------------|---------|--------------|
| input_4 (InputLayer) | [None, 78] | 0 | |
| input_5 (InputLayer) | [None, 57] | 0 | |
| embedding (Embedding) [0] | (None, 78, 300) | 5399100 | input_4[0] |
| embedding_1 (Embedding) [0] | (None, 57, 300) | 9022800 | input_5[0] |

In [36]:

```

1 import tensorflow_addons as tfa
2 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
3 import tensorflow as tf
4 # Learning rate schedule
5 def step_decay(epoch):
6
7     import math
8     initial_lrate = 0.1
9     drop = 0.9
10    epochs_drop = 10.0
11    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
12    return lrate
13
14 #lrate = LearningRateScheduler(step_decay)
15 lrate = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1, patience=10, m
16 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 10)
17 fname = 'best_model8.hdf5'
18 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save_
19 callbacks = [lrate,es]
20 #optimizer=tfa.optimizers.SGDW(weight_decay=0.99,momentum=0.9, Learning_rate=
21 #optimizer=tfa.optimizers.AdamW(weight_decay=0.99, Learning_rate=0.001)
22 optimizer=tf.optimizers.Adam(learning_rate=0.001)
23
24
25 #sgd = SGD(Learning_rate = 0.1, )
26 model8.compile(optimizer = optimizer, loss = 'categorical_crossentropy', met
27 model8.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 100, batc
28             validation_data = ([val_padded_1, val_padded_2], y_val_enc), call

```

Epoch 1/100
8584/8584 [=====] - 335s 38ms/step - loss: 0.8649 -
accuracy: 0.6084 - val_loss: 0.7642 - val_accuracy: 0.6739
Epoch 2/100
8584/8584 [=====] - 337s 39ms/step - loss: 0.7751 -
accuracy: 0.6686 - val_loss: 0.7407 - val_accuracy: 0.6846
Epoch 3/100
8584/8584 [=====] - 337s 39ms/step - loss: 0.7234 -
accuracy: 0.7017 - val_loss: 0.6674 - val_accuracy: 0.7303
Epoch 4/100
8584/8584 [=====] - 338s 39ms/step - loss: 0.6876 -
accuracy: 0.7219 - val_loss: 0.6438 - val_accuracy: 0.7430
Epoch 5/100
8584/8584 [=====] - 341s 40ms/step - loss: 0.6680 -
accuracy: 0.7338 - val_loss: 0.6266 - val_accuracy: 0.7544
Epoch 6/100
8584/8584 [=====] - 337s 39ms/step - loss: 0.6515 -
accuracy: 0.7426 - val_loss: 0.6153 - val_accuracy: 0.7569
Epoch 7/100
8584/8584 [=====] - 337s 39ms/step - loss: 0.6445 -
accuracy: 0.7485 - val_loss: 0.6125 - val_accuracy: 0.7595

In []:

1

In []:

1

tuning Bi-lstm, attention with glove

In [37]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length=
6                         weights = [embedding_matrix_glove_1], trainable = False)(in
7 #lstm_s1 = Bidirectional(LSTM(256, return_sequences = True) , merge_mode =
8 lstm1 = Bidirectional(LSTM(256, return_sequences = True, dropout=0.3, kernel_
9                         merge_mode = 'concat')(emb1)
10 att_1 = Attention_Custom(512, lstm1.shape)
11 att_weights_1, context_1 = att_1.call(lstm1)
12 s1_dense = Dense(units = 300, activation = 'relu' )(context_1)
13
14
15 #getting sentence2 representations
16 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
17 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length=
18                         weights = [embedding_matrix_glove_2], trainable = False)(in
19 #lstm_s2 = Bidirectional(LSTM(20, return_sequences = True))(emb2)
20 lstm2 = Bidirectional(LSTM(256, return_sequences = True, dropout=0.3, kernel_
21                         merge_mode = 'concat')(emb2)
22 att_2 = Attention_Custom(512, lstm2.shape)
23 att_weights_2, context_2 = att_2.call(lstm2)
24 s2_dense = Dense(units = 300, activation = 'relu' )(context_2)
25
26 #merging sentence1 and sentence2 representations
27 merged_representations = concatenate([s1_dense, s2_dense])
28 #point-wise product of sentence representations
29 prod = tf.math.multiply(s1_dense, s2_dense)
30 #point-wise absolute difference of sentence representations
31 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
32 abs_diff = tf.math.abs(subtracted)
33
34 #concatenating
35 concatenated = concatenate([merged_representations, prod, abs_diff])
36 # concatenated = tf.expand_dims(concatenated, axis=-1)
37 # conv1 = Conv1D(filters = 64, kernel_size = 3, strides=1, padding='valid')(
38 # conv2 = Conv1D(filters = 32, kernel_size = 3, strides=1, padding='valid')(
39 # conv3 = Conv1D(filters = 16, kernel_size = 3, strides=1, padding='valid')(
40
41 # flat = Flatten()(conv3)
42 #MLP and three way classifier
43 dense1 = Dense(units = 512, activation = 'relu')(concatenated)
44 dense2 = Dense(units = 128, activation = 'relu')(dense1)
45 #drop1 = Dropout(0.3)(dense1)
46 # dense2 = Dense(units = 50, activation = 'relu')(drop1)
47 output = Dense(units = 3, activation = 'softmax')(dense2)
48 model9 = Model(inputs = [input1, input2], outputs = output)
49
50

```

```
In [39]: 1 model9.summary()
```

Model: "model_2"

| Layer (type) | Output Shape | Param # | Connected to |
|--------------------------------|-----------------|---------|--------------|
| input_6 (InputLayer) | [None, 78] | 0 | |
| input_7 (InputLayer) | [None, 57] | 0 | |
| embedding_2 (Embedding) [0] | (None, 78, 300) | 5399100 | input_6[0] |
| embedding_3 (Embedding) [0] | (None, 57, 300) | 9022800 | input_7[0] |

In [40]:

```

1 import tensorflow_addons as tfa
2 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
3 import tensorflow as tf
4 lrate = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1, patience=10, m
5 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 5)
6 fname = 'best_model9.hdf5'
7 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save
8 callbacks = [lrate,es]
9 #optimizer=tfa.optimizers.SGDW(weight_decay=0.99,momentum=0.9,learning_rate=
10 #optimizer=tfa.optimizers.AdamW(weight_decay=0.99,learning_rate=0.001)
11 optimizer=tf.optimizers.Adam(learning_rate=0.001)
12
13
14 #sgd = SGD(Learning_rate = 0.1, )
15 model9.compile(optimizer = optimizer, loss = 'categorical_crossentropy', met
16 model9.fit([train_padded_1, train_padded_2], y_train_enc, epochs = 100, batc
17 validation_data = ([val_padded_1, val_padded_2], y_val_enc), call

```

Epoch 1/100
8584/8584 [=====] - 378s 43ms/step - loss: 0.7961 - accuracy: 0.6583 - val_loss: 0.6676 - val_accuracy: 0.7347
Epoch 2/100
8584/8584 [=====] - 379s 44ms/step - loss: 0.6795 - accuracy: 0.7322 - val_loss: 0.6228 - val_accuracy: 0.7649
Epoch 3/100
8584/8584 [=====] - 378s 44ms/step - loss: 0.6436 - accuracy: 0.7542 - val_loss: 0.6054 - val_accuracy: 0.7734
Epoch 4/100
8584/8584 [=====] - 381s 44ms/step - loss: 0.6209 - accuracy: 0.7679 - val_loss: 0.5962 - val_accuracy: 0.7821
Epoch 5/100
8584/8584 [=====] - 379s 44ms/step - loss: 0.6051 - accuracy: 0.7781 - val_loss: 0.5895 - val_accuracy: 0.7888
Epoch 6/100
8584/8584 [=====] - 381s 44ms/step - loss: 0.5924 - accuracy: 0.7869 - val_loss: 0.5746 - val_accuracy: 0.7937
Epoch 7/100
8584/8584 [=====] - 388s 45ms/step - loss: 0.5822 - accuracy: 0.7930 - val_loss: 0.5774 - val_accuracy: 0.7972
Epoch 8/100
8584/8584 [=====] - 387s 45ms/step - loss: 0.5724 - accuracy: 0.7991 - val_loss: 0.5743 - val_accuracy: 0.7989
Epoch 9/100
8584/8584 [=====] - 383s 45ms/step - loss: 0.5651 - accuracy: 0.8040 - val_loss: 0.5880 - val_accuracy: 0.7973
Epoch 10/100
8584/8584 [=====] - 383s 45ms/step - loss: 0.5583 - accuracy: 0.8083 - val_loss: 0.5851 - val_accuracy: 0.7995
Epoch 11/100
8584/8584 [=====] - 387s 45ms/step - loss: 0.5527 - accuracy: 0.8126 - val_loss: 0.5864 - val_accuracy: 0.8010
Epoch 12/100
8584/8584 [=====] - 387s 45ms/step - loss: 0.5465 - accuracy: 0.8157 - val_loss: 0.5822 - val_accuracy: 0.8029
Epoch 13/100
8584/8584 [=====] - 390s 45ms/step - loss: 0.5419 - accuracy: 0.8187 - val_loss: 0.5809 - val_accuracy: 0.8058

```
Epoch 14/100
8584/8584 [=====] - 388s 45ms/step - loss: 0.5365 - accuracy: 0.8221 - val_loss: 0.5994 - val_accuracy: 0.8005
Epoch 15/100
8584/8584 [=====] - 389s 45ms/step - loss: 0.5328 - accuracy: 0.8250 - val_loss: 0.6054 - val_accuracy: 0.7972
Epoch 16/100
8584/8584 [=====] - 389s 45ms/step - loss: 0.5286 - accuracy: 0.8275 - val_loss: 0.6007 - val_accuracy: 0.8097
Epoch 17/100
8584/8584 [=====] - 388s 45ms/step - loss: 0.5257 - accuracy: 0.8293 - val_loss: 0.6020 - val_accuracy: 0.8063
Epoch 18/100
8584/8584 [=====] - 375s 44ms/step - loss: 0.5208 - accuracy: 0.8318 - val_loss: 0.6119 - val_accuracy: 0.8055
Epoch 19/100
8584/8584 [=====] - 381s 44ms/step - loss: 0.5174 - accuracy: 0.8346 - val_loss: 0.6115 - val_accuracy: 0.8087
Epoch 20/100
8584/8584 [=====] - 381s 44ms/step - loss: 0.5147 - accuracy: 0.8360 - val_loss: 0.6146 - val_accuracy: 0.8100
Epoch 21/100
8584/8584 [=====] - 381s 44ms/step - loss: 0.5119 - accuracy: 0.8385 - val_loss: 0.6076 - val_accuracy: 0.8078
Epoch 22/100
8584/8584 [=====] - 378s 44ms/step - loss: 0.5088 - accuracy: 0.8397 - val_loss: 0.6135 - val_accuracy: 0.8062
Epoch 23/100
8584/8584 [=====] - 587s 68ms/step - loss: 0.5061 - accuracy: 0.8418 - val_loss: 0.6267 - val_accuracy: 0.8039
Epoch 24/100
8584/8584 [=====] - 384s 45ms/step - loss: 0.5033 - accuracy: 0.8431 - val_loss: 0.6533 - val_accuracy: 0.8062
Epoch 25/100
8584/8584 [=====] - 385s 45ms/step - loss: 0.5006 - accuracy: 0.8447 - val_loss: 0.6251 - val_accuracy: 0.8045
```

Out[40]: <keras.callbacks.History at 0x216243f4b38>

In []: 1

Bi-LSTM with Attention on fasttext embeddings and extracted features

In [49]:

```

1 emb_dim = 300
2
3 #getting sentence 1 representations
4 input1 = Input(shape=(maxlen1,)) #for sentence1 embeddings
5 emb1 = Embedding(input_dim = vocab_size_1, output_dim = emb_dim, input_length =
6 weights = [embedding_matrix_glove_1], trainable = False)(in
7 #lstm_s1 = Bidirectional(LSTM(256, return_sequences = True) , merge_mode =
8 lstm1 = Bidirectional(LSTM(256, return_sequences = True, dropout=0.3, kernel_
9 merge_mode = 'concat')(emb1)
10 att_1 = Attention_Custom(512, lstm1.shape)
11 att_weights_1, context_1 = att_1.call(lstm1)
12 print(context_1.shape)
13 s1_dense = Dense(units = 300, activation = 'relu' )(context_1)
14 print(s1_dense.shape)
15
16 #getting sentence2 representations
17 input2 = Input(shape=(maxlen2,)) #for sentence2 embeddings
18 emb2 = Embedding(input_dim = vocab_size_2, output_dim = emb_dim, input_length =
19 weights = [embedding_matrix_glove_2], trainable = False)(in
20 #lstm_s2 = Bidirectional(LSTM(20, return_sequences = True))(emb2)
21 lstm2 = Bidirectional(LSTM(256, return_sequences = True, dropout=0.3, kernel_
22 merge_mode = 'concat')(emb2)
23 att_2 = Attention_Custom(512, lstm2.shape)
24 att_weights_2, context_2 = att_2.call(lstm2)
25 s2_dense = Dense(units = 300, activation = 'relu' )(context_2)
26 print(s2_dense.shape)
27 #merging sentence1 and sentence2 representations
28 merged_representations = concatenate([s1_dense, s2_dense])
29 #point-wise product of sentence representations
30 prod = tf.math.multiply(s1_dense, s2_dense)
31 #point-wise absolute difference of sentence representations
32 subtracted = tf.keras.layers.Subtract()([s1_dense, s2_dense])
33 abs_diff = tf.math.abs(subtracted)
34
35 input3 = Input(shape=(29,)) #for extracted features
36 dense_1 = Dense(units = 16, activation = 'relu')(input3)
37
38 #concatenating
39 concatenated = concatenate([merged_representations, prod, abs_diff, dense_1])
# concatenated = tf.expand_dims(concatenated, axis=-1)
41 # conv1 = Conv1D(filters = 64, kernel_size = 3, strides=1, padding='valid')
42 # conv2 = Conv1D(filters = 32, kernel_size = 3, strides=1, padding='valid')
43 # conv3 = Conv1D(filters = 16, kernel_size = 3, strides=1, padding='valid')
44
45 # flat = Flatten()(conv3)
46 #MLP and threeway classifier
47 dense1 = Dense(units = 512, activation = 'relu')(concatenated)
48 dense2 = Dense(units = 128, activation = 'relu')(dense1)
49 #drop1 = Dropout(0.3)(dense1)
50 # dense2 = Dense(units = 50, activation = 'relu')(drop1)
51 output = Dense(units = 3, activation = 'softmax')(dense2)
52 model10 = Model(inputs = [input1, input2, input3], outputs = output)
53
54

```

(None, 512)
 (None, 300)

(None, 300)

In [26]: 1 model10.summary()

Model: "model"

| Layer (type) | Output Shape | Param # | Connected to |
|--------------------------------|-----------------|---------|--------------|
| input_2 (InputLayer) | [None, 78] | 0 | |
| input_3 (InputLayer) | [None, 57] | 0 | |
| embedding_1 (Embedding) [0] | (None, 78, 300) | 5399100 | input_2[0] |
| embedding_2 (Embedding) [0] | (None, 57, 300) | 9022800 | input_3[0] |

In [27]:

```

1 import tensorflow_addons as tfa
2 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping,
3 import tensorflow as tf
4 lrate = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1, patience=10, m
5 es = EarlyStopping(monitor='val_accuracy', mode='max', patience = 5)
6 fname = 'best_model10.hdf5'
7 checkpoint = ModelCheckpoint(fname, monitor="val_accuracy", mode="max", save
8 callbacks = [lrate,es]
9 #optimizer=tfa.optimizers.SGDW(weight_decay=0.99,momentum=0.9,learning_rate=
10 #optimizer=tfa.optimizers.AdamW(weight_decay=0.99,learning_rate=0.001)
11 optimizer=tf.optimizers.Adam(learning_rate=0.001)
12
13
14 #sgd = SGD(Learning_rate = 0.1, )
15 model10.compile(optimizer = optimizer, loss = 'categorical_crossentropy', me
16 model10.fit([train_padded_1, train_padded_2, X_train_1], y_train_enc, epochs
17 validation_data = ([val_padded_1, val_padded_2, X_val_1], y_val_e

```

Epoch 1/100
8584/8584 [=====] - 386s 44ms/step - loss: 0.8149 - accuracy: 0.6518 - val_loss: 0.6799 - val_accuracy: 0.7253
Epoch 2/100
8584/8584 [=====] - 393s 46ms/step - loss: 0.7116 - accuracy: 0.7116 - val_loss: 0.6541 - val_accuracy: 0.7416
Epoch 3/100
8584/8584 [=====] - 393s 46ms/step - loss: 0.6529 - accuracy: 0.7451 - val_loss: 0.6061 - val_accuracy: 0.7687
Epoch 4/100
8584/8584 [=====] - 393s 46ms/step - loss: 0.6226 - accuracy: 0.7624 - val_loss: 0.5893 - val_accuracy: 0.7775
Epoch 5/100
8584/8584 [=====] - 394s 46ms/step - loss: 0.6039 - accuracy: 0.7732 - val_loss: 0.5804 - val_accuracy: 0.7899
Epoch 6/100
8584/8584 [=====] - 394s 46ms/step - loss: 0.5888 - accuracy: 0.7826 - val_loss: 0.5782 - val_accuracy: 0.7843
Epoch 7/100
8584/8584 [=====] - 396s 46ms/step - loss: 0.5801 - accuracy: 0.7881 - val_loss: 0.5885 - val_accuracy: 0.7877
Epoch 8/100
8584/8584 [=====] - 391s 46ms/step - loss: 0.5766 - accuracy: 0.7918 - val_loss: 0.5671 - val_accuracy: 0.7991
Epoch 9/100
8584/8584 [=====] - 392s 46ms/step - loss: 0.5669 - accuracy: 0.7968 - val_loss: 0.5760 - val_accuracy: 0.7914
Epoch 10/100
8584/8584 [=====] - 391s 46ms/step - loss: 0.5607 - accuracy: 0.8010 - val_loss: 0.5740 - val_accuracy: 0.7993
Epoch 11/100
8584/8584 [=====] - 391s 46ms/step - loss: 0.5560 - accuracy: 0.8045 - val_loss: 0.5765 - val_accuracy: 0.7980
Epoch 12/100
8584/8584 [=====] - 386s 45ms/step - loss: 0.5485 - accuracy: 0.8083 - val_loss: 0.5911 - val_accuracy: 0.7930
Epoch 13/100
8584/8584 [=====] - 384s 45ms/step - loss: 0.5427 - accuracy: 0.8112 - val_loss: 0.5747 - val_accuracy: 0.7977

```
Epoch 14/100
8584/8584 [=====] - 384s 45ms/step - loss: 0.5374 - accuracy: 0.8147 - val_loss: 0.5971 - val_accuracy: 0.7932
Epoch 15/100
8584/8584 [=====] - 384s 45ms/step - loss: 0.5313 - accuracy: 0.8179 - val_loss: 0.5945 - val_accuracy: 0.7986
```

Out[27]: <keras.callbacks.History at 0x191810855f8>

the best performance by a model until now is val acc = 0.81, model: bi-lstm with self attention and glove embeddings

In []: 1